

To Facilitate or not to Facilitate: Human and LLM Facilitator Tendencies in Online Discussions

Anonymous ACL submission

Abstract

Automating facilitation in online discussions is a long-standing social concern given the increasing time we spend on online spaces and the failure of content moderation approaches. While studies have been conducted on how to facilitate, none have answered the essential question of when to do so. A potential answer is using LLMs, which ostensibly make automated, large-scale intervention increasingly feasible. In this study, we examine when LLMs decide to facilitate by defining what facilitation is, observing when humans decide to facilitate, and comparing their decisions with those made by LLMs. To this end, we create PEFK, a corpus standardizing and aggregating all relevant facilitation datasets. We are the first to run a survey on facilitation timing, which we execute using expert facilitative participants and LLM-as-a-judge models. We discover that while humans are more cautious, LLMs are excessively eager to facilitate, although both are more certain when judging that facilitation is not needed. We then investigate whether this behavior can be corrected using alternative setups for LLMs and training ModernBert classifiers on established datasets, finding that the latter perform more reliably than the former, although current datasets result in relatively low performance.

1 Introduction

Given the scale of social media networks and online spaces, facilitation needs to be automated to keep up with the amount of posted content (Small et al., 2023; Burton et al., 2024). So far, platforms have largely ignored the issue, preferring automated flag-and-remove mechanisms (“content moderation”) (Korre et al., 2025), which are not sufficient (Cresci et al., 2022; Trujillo and Cresci, 2022). While recent work has undertaken the task of choosing *how* to facilitate (Chen et al., 2025; Schroeder et al., 2024; Park et al., 2012), there has been little interest in *when* to do so (Falk et al., 2021).




Figure 1: At any point in a discussion, the facilitator must decide whether to intervene or not. Not intervening when the situation calls for it may lead to the discussion worsening (e.g., topic derailment, escalation, missed participation). On the other hand, intervening when the discussion is already going well may lead to irritation among participants.

Deciding when to intervene in a discussion is a matter of great importance. Intervening too early or too frequently may lead to users expressing frustration at being constantly interrupted, while intervening too late may lead to a significant escalation and derailment of the discussion (Fig. 1, Korre et al. (2025); Gao et al. (2025)). Due to the sheer scale of modern online spaces automation is needed, which the field of Natural Language Processing (NLP) has attempted to implement using models to determine when an intervention is necessary. Encoder-Only (EncOn) classifiers were used to this end (Falk et al., 2021), although recently LLMs have been championed as the eventual solution to online facilitation (Korre et al., 2025; Small et al., 2023). However, previous work in LLM facilitation indicates that LLM facilitators are way too eager to intervene, rendering them unusable as autonomous facilitative agents (Tsirmpas et al., 2026). Yet, this finding has not been contrasted with human tendencies. In this study, we focus on the following Research Question (RQ):

RQ. *How and why do LLMs differ with humans when deciding when to facilitate?*

In order to answer the central research question we need to answer the following sub-questions:

What constitutes a facilitative intervention?

There is currently no consensus on how to define a facilitative intervention. By examining relevant literature, we find three common definitions implicitly used in prior work; *escalation*, *professional facilitation*, and *facilitative comments* (§2.1).

Where do humans facilitate? While the majority of content where facilitation is needed is posted on online, written discussions, recent work predominantly focuses on *domain-specific oral discussion* scenarios (e.g., radio debates). Since datasets for discussion facilitation are generally few, small, and straddle radically different domains (e.g., debates, forums, deliberations), we gather all datasets relating to any of the three possible definitions of the term and compile the “Prosocial and Effective Facilitation in Konversations (PEFK)”, a standardized dataset for discussion facilitation with a unified data and label schema (§3).

How do humans facilitate? We create and execute a survey task designed to gauge human facilitative patterns by sampling 1226 discussion chunks from our dataset, and asking 10 expert facilitative participants whether they would intervene as facilitators. We find that human participants tend to take a cautious approach to facilitation, generally preferring not to intervene and expressing greater confidence in not doing so (Finding 1). Additionally, human participants may pay attention to different spans of a discussion, which can lead them to choose different courses of action regarding facilitation.

How does LLM facilitation differ? We repeat the survey with six open-source LLMs. Unlike humans, LLMs tend to over-intervene, confirming previous studies (Tsirmpas et al., 2026) and rely more heavily on positive reinforcement. However, compared to humans, they may be more consistent between themselves, hinting that they are likely focusing on similar spans of the discussion. We also find that, much like humans, determining when not to facilitate is easy—the opposite is not (Findings 2,3).

How can we fix LLM facilitation? We investigate whether aggregating all currently available

datasets can support either the training of classifiers or the use of out-of-the-box LLMs for a simpler and more commonly studied task in the literature: predicting whether a facilitator would intervene, rather than whether a facilitative comment is needed (§5). Our results show that LLMs are less reliable than ModernBert (Warner et al., 2025) classifiers, although their performance depends heavily on the model used (Finding 4).

Overall, our study offers the following contributions: We bridge differences between prior work on (i) what a facilitation intervention is, (ii) oral and written discussions (1), by compiling the largest standardized facilitation dataset so far (2). Through the execution of an extensive survey performed by experts on facilitation (3), we show that the key difference between humans and LLMs rests in how frequently they deem facilitation to be necessary (4). We are the first to evaluate LLMs for facilitation timing prediction, finding that they are wildly inconsistent while EncOn models offer a more reliable alternative, though both approaches remain constrained by the limited current understanding of facilitation within NLP (5).

Ultimately, achieving reliable automated facilitation requires moving beyond surface signals and using comprehensive, human-informed data that captures the true complexity of human facilitation. We release the code for the creation of our dataset and the experiments presented in this paper under a GPLv3-or-later license.¹

2 Related Work and Background

2.1 What are facilitative interventions?

Facilitation (or conversational moderation) refers to the active participation of people in order to improve discussions (Korre et al., 2025). Relevant work assumes that facilitation is conducted by specially trained staff (Schroeder et al., 2024; Park et al., 2012; Chen et al., 2025), or users appointed by the community (Schaffner et al., 2024). Studying relevant literature on facilitation yields multiple definitions for what a facilitative intervention is, which can be categorized to three main definitions:

Definition 1 (Escalation). A facilitator needs to intervene at the precise point at which, without facilitation, the discussion would derail or escalate.

¹Dataset: <https://anonymous.4open.science/r/facilitation-dataset-37DF>, Experiments: <https://anonymous.4open.science/r/interventions-6053>, Survey Data: <https://anonymous.4open.science/r/survey-data-C43E>

Escalation is usually measured as the point where the participants request arbitration (e.g., escalate the matter to a site moderator in Wikipedia (De Kock and Vlachos, 2021; Chang and Danescu-Niculescu-Mizil, 2019)). This definition is quite appealing for NLP researchers as the labels are easy to acquire since a ground truth exists, and the existence of a single, largely unambiguous ground-truth is ideal for the training and evaluation of any Machine Learning (ML)-based classifier. However, a facilitator has multiple duties that do not directly relate to preventing escalation such as summarizing discussions (Small et al., 2023), managing participation (Chen et al., 2025; Schroeder et al., 2024), and providing important information for new members (Park et al., 2012). Furthermore, an unambiguous signal that a discussion has escalated is not often provided by platforms (Table 2). Thus, this definition is not considered in this study, but is included in our dataset (§3).

Definition 2 (Professional facilitation). Facilitation should occur at the point where a professional facilitator would have intervened.

In current literature, it is generally accepted that all comments made by a facilitator should be considered facilitative interventions; for instance, only these comments are assumed to follow some facilitative strategy (Falk et al., 2021; Chen et al., 2025; Schroeder et al., 2024). Defining facilitative comments this way is simple and intuitive, but can only be used in domains where (1) professional facilitators exist and (2) they can be identified. These conditions exclude the vast majority of online discussions, as very few web-crawled datasets include such information.²

Definition 3 (Human judgment). Facilitation should only occur when the majority agree that it is necessary.

There is an argument to be made that professional facilitators can make more accurate judgments due to expertise. However, in reality most facilitation is handled by either volunteers elevated to the position by their community (Schaffner et al., 2024) or by the users themselves in an ad-hoc fashion (Falk et al., 2024). In this case, the task is no longer about detecting or predicting whether *a facilitator would speak*, but rather the much more difficult and subjective task of whether the comment is *facilitative in nature*.

²See Table 2, where the largest datasets such as WikiConv (Hua et al., 2018) do not contain such information.

Name	Text	D1	D2	D3
Joe	Thank you for sharing. So I am Joe from Chicago originally...	×	×	×
Marie	Thank you Joe.	×	✓	?
Peter	Yeah. So I'll kind of go through it again. I am Peter...	×	×	×
Marie	Thank you, Peter. So I think we have another participant on the line...	×	✓	✓

Table 1: Example text where facilitation interventions are identified according to each definition (D1-3). *Marie* is a facilitator. We assume the discussion does not derail. The second comment *may* be facilitative according to Def. 3, which depends on the consensus of participants.

What is the difference? Table 1 shows an example where the same discussion is annotated following the three definitions. Def. 1 is reliant on the outcome of the discussion, and would miss examples where successful interventions led to a smooth discussion. Def. 2 would give more fine-grained information, although we run the risk of classifying non-facilitative comments as facilitative (and sometimes the opposite (Falk et al., 2024)). Finally, Def. 3 can provide us with fine-grained information, although ambiguous cases remain. It is worth noting that these definitions are not mutually exclusive, as demonstrated in Table 1.

2.2 Determining when to facilitate

Using models The only work to our knowledge that has explicitly attempted to predict facilitative interventions (Falk et al., 2021) uses *EncOn* classifiers on a single dataset—Cornell eRulemaking Initiative (CeRI) (Park et al., 2012)—with modest performance. Indeed, these classifiers have proven competent in the context of escalation prediction (Chang and Danescu-Niculescu-Mizil, 2019; Yuan and Singh, 2023), and facilitative tactics detection (Chen et al., 2025; Schroeder et al., 2024). Recent work (Schroeder et al., 2024; Chen et al., 2025) uses LLM-as-a-judge models to annotate data in the field of facilitation, although recent research suggests that fully-automated facilitation is still out of our reach (Korre et al., 2025).

Using humans Expert human facilitators are frequently used in research (Chen et al., 2025; Park et al., 2012; Schroeder et al., 2024), but this is not a scalable process due to costs (Rossi et al., 2024).

Ignoring it Most platforms hold users responsible for maintaining the quality of discussion and the

enforcement of community and platform rules, especially in jurisdictions where the platforms themselves are largely absolved of the responsibility of posted content, such as the U.S. (Seering, 2020; Schaffner et al., 2024). In some cases, the users are responsible for just flagging and reporting content (Seering, 2020), while platforms like Reddit use volunteer moderators, who also function as community leaders and facilitators (Seering, 2020; Falk et al., 2024), raising ethical questions about “Moderation as Free Labor” (Matias, 2019).

3 PEFK: A large, multi-domain facilitation dataset

Since there is currently no dataset that explicitly concerns itself with facilitation interventions, we construct the “Prosocial and Effective Facilitation in Konversations (PEFK)” dataset, a standardized, multi-domain compilation of the facilitation-related datasets listed by Korre et al. (2025) (see Table 2). PEFK contains data from both online forums and oral discussions, and standardizes metadata, metrics, and labeling schemes according to a simple, unified schema, making it suitable as a benchmark on facilitator interventions. More information on preprocessing steps and dataset schema can be found in App. B.

PEFK is designed to facilitate exploration and analysis of facilitative interventions under all three definitions described in §2.1. Specifically, we use a subset of the data that includes professional facilitation annotations to train our models according to Def. 2 in §5 (see App. B.3; Fig. 5). We also use a subset of PEFK to survey human facilitative patterns according to Def. 3. Finally, we include escalation and toxicity, as used in Def. 1, although this information is not utilized in this study (see §2.1).

4 Do humans and LLMs intervene at similar circumstances?

To support a systematic investigation of facilitative behavior under Def. 3, we conduct a survey on PEFK to analyze facilitation timing patterns between humans and LLMs. Rather than establishing a ground truth through Inter-Annotator Agreement (IAA), we adopt a survey-based approach, since prior work has shown that obtaining reliable annotations for similar tasks is highly challenging, particularly in newly explored settings such as ours (see “Survey Procedure”).

4.1 Facilitative survey

Labels To avoid turning facilitation into a moderation task (i.e., using reactive responses after unwanted comments/utterances), we introduce two types of interventions: positive and negative reinforcement. *Positive reinforcement*, aims to encourage or reward constructive behavior (e.g., expressing appreciation or affirming contributions), while *negative reinforcement*, aims to discourage undesirable behavior or guide interlocutors back to productive discourse (e.g., addressing incivility or off-topic comments). This distinction represents a tool through which we ensured the survey participants did not default to content moderation, which we frequently observed during discussions with them, and during pilot survey rounds—it does not represent a theoretical contribution. Nevertheless, it provides an extra signal which we utilize in our analysis.³ We adopt an ordinal scale approach, where the survey participants rate each dialogue for either type of reinforcement (positive or negative), as well as the possibility of no reinforcement using a scale from 1 to 5.⁴

Data We randomly split multi-party discussions from the PEFK dataset into small chunks containing a set amount of comments (6 comments per discussion), and select a random sample for the survey (1224 discussions—721 remained after the participants flagged some excerpts as malformed). This design aims to replicate a known limitation in online asynchronous facilitation scenarios; the facilitator is expected to decide on a course of action considering only the last parts of a discussion, as shown by several studies (Hewitt, 2003; Park et al., 2016; Kümpel and Unkel, 2021). In the case of oral discussions, there is evidence to suggest that the most recent comments are given relatively greater importance (Crano, 1977; Cockburn et al., 2017).

Survey Participants We recruited 11 experts on Computational Social Science with extensive experience with NLP tasks from sentiment analysis to persuasion strategy detection. We preferred expert

³There is some theoretical grounding for using this approach is described in App. D.2. Similarly, Falk et al. (2021) use high and low Argument Quality for further analysis of facilitation prediction, although, unlike their approach, we obtain the explanatory signal directly from the survey participants, rather than established EncOn models.

⁴In order to transform the multiple ordinal labels into an easily manageable form for the purposes of analysis, we transform them into a single multiclass label schema as described in App. D.3.

Name	#Comments	#Discussions	Domain	Format	Fac.
WikiDisputes (2021)	96,320	8,727	Forum	Text	×
WikiTactics (WT) (2022)	3,850	213	Forum	Text	✓
WikiConv (2018)	17,806,373	4,486,983	Forum	Text	×
Conversations Gone Awry / CMV II (2019)	40,607	6,691	Forum	Text	×
Cornell eRulemaking Initiative (CeRI) (2012)	3,700	132	Forum	Text	✓
User Moderation (UMOD) (2024)	2,000	1,000	Forum	Text	✓
Intelligence Squared (IQ2) (2016)	34,245	108	Formal debate	Oral	✓
Why-How-Who (WHoW) (2025)	25,542	178	Radio / TV	Oral	✓
Fora (2024)	39,438	262	Story sharing	Oral	✓

Table 2: The datasets used in PEFK include a large variety of domains, some crawled, and others as transcripts from oral discussions. All datasets include multi-participant discussions (besides UMOD which contains comment pairs). The **Fac.** column denotes whether information on which users are facilitators is available.

instead of crowdsourced participants as the task is highly technical, subjective, and mentally demanding. Additionally, crowdsourced participants have been shown to exhibit more bias in general [CITE ME], which concerned us since we tried to avoid confusion with content moderation—a better known task. Our concerns were validated when preliminary discussions with participants revealed that they indeed considered their task to be content moderation.

The gender distribution was relatively balanced with 4 males and 7 females, while their ages had a median of 30 years. Before the commencement of the survey, the participants completed an agreement form in compliance with GDPR (European Union, 2016) ethical requirements and were compensated in accordance with national legislation. More information about the participant selection can be found in App. E.

Survey Procedure Unlike prior work that sought to establish consensus among facilitators (Falk et al., 2024; Schroeder et al., 2024; Chen et al., 2025), our goal is instead to explore facilitation patterns without presuming that such a consensus exists. Indeed, facilitation is a highly subjective task, and substantial disagreements persist even among experienced professionals operating within the same communities and under the same rule-sets (Koshy et al., 2024); consensus remains elusive even in cases where the task is less ambiguous and the dataset more focused (Schroeder et al., 2024).⁵

Before the survey, the participants were given a training task where they categorized facilitative tactics according to established taxonomies on PEFK (see App. D.1) and given ample resources and time

⁵During the survey, we inserted duplicated entries in secret to estimate consistency, finding that all-but-one participants remained 100% consistent in their judgments during the task—see App. D.5.

to familiarize themselves with online facilitation. We felt that training was a necessary step in order to disambiguate facilitation from content moderation after preliminary discussions with the participants. While it could in principle bias the participants, we attempted to control potential bias by strictly limiting the training scope to strategies only; there was no implicit or explicit input regarding the timing or flow of conversation (see App. D).

The main survey task required the participants to rate instances from PEFK by determining whether, at any point within a specific discussion chunk, reinforcement was needed, as well as what kind. The participants could optionally add rationales with which they could justify their rating. The full survey guidelines can be found in App. G. We ran two pilot surveys to clear up misunderstandings around the task and incorporate participant feedback, before executing the main task. Additionally, we conducted one-on-one interviews and weekly group meetings concurrently with the main task. More details can be found in App. E.

LLM Survey We use a variety of open-source LLMs (App. C.1) to ensure reproducibility. The prompts used were adopted from the guidelines provided to the experts—see App. H.

4.2 Human facilitation

Humans make cautious facilitators The main dilemma around facilitative timing is determining whether the facilitator prioritizes preventing potentially harmful speech, or letting the conversation continue naturally (Fig. 1). According to Fig. 2, most participants generally remain conservative—most choose not to intervene in a large majority of cases, although some (e.g., A3, A8, and A9) exhibit comparatively more proactive behav-




Figure 2: Distribution of all classes per participant. Human participants are either very cautious (A2, A4–A7, A10, A11) or somewhat cautious (A3, A8, A9). On the other hand, the LLMs (b) are significantly more eager to facilitate than humans (a).




Figure 3: Certainty of judgments per facilitation label. LLMs (orange bars on the right) are much more confident in their judgments than humans (high frequencies on the extreme modes i.e., 1 and 5).

ior with higher rates of positive interventions.⁶ Positive-intervention rates also appear to cluster around two broad ranges (approximately 20–30% and 35–40%), which is consistent with the tendencies of professional facilitators observed in the original dataset (see Fig. 5; App. B.2).⁷

Human facilitators may attend over different discussion spans To qualitatively understand the sources of divergence in human facilitative patterns, we analyzed a “worst-case” scenario: situations where half of the participants favored positive reinforcement and the other half favored negative reinforcement. As shown in Table 3, a specific discussion example illustrates this conflict. Participants justifying negative reinforcement cited the conver-

⁶One participant (A1) did not complete the task and was therefore not considered for the analysis.

⁷These observations are consistent across datasets—see App. D.4—reinforcing our motivation of transferring observations from oral datasets to online, written facilitation.

#	Content
1	Well nothing is really *inherently* bad...
2	Well political correctness is about avoiding offense...
4	Are you really arguing that WWI was unjust towards cisgendered men?
5	Not WWI itself, but that period, and I don't see what is so odd about it...
6	There were transgendered people in WWI. You still haven't justified how it was genuinely an injustice...

Table 3: Sections from the same discussion chunk. Participants facilitate differently depending on which section they decide is more important.

	Qwen32b	Qwen7b	Llama8b	Annotator	Olmo7b	Olmo32b	Llama70b	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11
706	Pos	Both	Pos	Pos	Pos	Pos	Pos	Neg	No	No	No	No	No	No	No	No	No
582	Pos	Both	Pos	Pos	Pos	Pos	Pos	Neg	No	No	Neg	Neg	No	No	No	Both	Neg
520	Pos	Both	Pos	Pos	Pos	Pos	Pos	No	Neg	No	No	No	No	No	No	No	No
459	Pos	Both	Pos	Pos	Pos	Pos	Pos	No	Neg	No	No	No	No	No	No	Both	No
661	Pos	Both	Pos	Pos	Pos	Pos	Pos	Neg	No	No	No	No	No	No	No	No	Neg

Figure 4: Top 5 discussion instances where the LLM majority vote disagrees with all human participants. The y-axis indicates the discussion instance index in the dataset. Labels follow the survey schema described in App. D.2: POS = Positive Reinforcement, NEG = Negative Reinforcement, BOTH = Both Positive and Negative Reinforcement, and NO = No Reinforcement.

sation’s trajectory toward escalation, suggesting intervention was necessary to curb aggressive language. Conversely, the same participants noted that certain comments—particularly the first two—were examples of politeness and hedging, making them candidates for positive reinforcement.

Finding 1. *Most human participants tend to be either highly cautious or to balance proactivity with caution. Disagreements may arise because participants focus on different spans of the discussion.*

4.3 Do LLMs facilitate differently?

The facilitative style of LLMs differs significantly Fig. 3 shows that humans are less certain than LLMs when it comes to positive reinforcement but in both cases the values are concentrated in the middle part of the distribution. For negative reinforcement, both humans and LLMs show high uncertainty, while for no reinforcement there is a clear difference with humans being more certain than LLMs.

Why exactly do LLMs differ from humans? We identify cases where LLMs largely agreed among themselves but completely disagreed with all humans in order to highlight situations where model behavior diverges acutely from human judgment. In total, there are 36 cases where no human agrees with the majority voting of the LLMs. Fig. 4

shows the top 5 extreme cases. Despite strong LLM consensus (agreement ≥ 0.66), all human participants labeled these comments differently.

Examining the five examples qualitatively, we see that all cases regard a narrative of the participants rather than a debate like discussion that could possibly lead to escalation.⁸ Specifically, discussion 706 is about gender, labor market change, and education, where multiple speakers exchange viewpoints on shifting economic roles.⁹ The excerpt contains a lot of anecdotal examples (e.g., references to one’s workplace examples and “Aviation High School”), rather than decision-oriented, focusing on broad social interpretation. In general, the dialogue is very self-contained, with speakers building perspectives rather than seeking resolution or guidance. This narrative like structure lead to the participant to choose to not to intervene and reporting that discussions like these comprise “neutral comments”. LLMs, on the other hand, being more prone to intervention, could interpret this neutrality as an opportunity for positive reinforcement. This is a hypothesis that should be verified with rationales by the participants in future research.

Finding 2. *LLMs are more intrusive and more consistent compared to humans, being eager to view neutral discussions as opportunities for positive reinforcement.*

5 Can we use existing datasets?

As discussed in §2.1, facilitation interventions can be defined in two ways; whether a comment is facilitative in nature (Def. 3), or whether a dedicated facilitator would participate at a given point in time (Def. 2). In this section, we investigate whether we can use available datasets to train and evaluate models using the latter definition. See App. F.1 for experiments focusing on the broader understanding of facilitation by LLMs and App. F.2 for ablation studies on both self-consistency and instruction-sensitivity.

Experimental setup Def. 2 and Def. 3 are radically different, and thus require completely sep-

⁸Unfortunately, we cannot rely on the LLMs for explanations post-hoc, nor use Chain of Thought (CoT) prompting to extract rationales, since both produce non-representative and deceptive explanations (Sanwal, 2025; Barez et al., 2025; Turpin et al., 2023). How the models arrived to these conclusions must therefore be evaluated from our own observations.

⁹Since each discussion span can involve as many as 30,000 characters, we refer the reader to the supplementary material for context.

Model	Pr.	Rec.	$F1_p$	$F1_n$	Sup.
CeRI					
LLaMa70B	0.152	0.082	0.107	0.783	335
LLaMa8B	0.246	0.518	0.333	0.567	335
ModBert	0.319	0.490	0.386	0.706	192
OLMo32B	0.154	0.071	0.097	0.795	335
OLMo7B	0.251	0.988	0.401	0.000	335
Qwen32B	0.100	0.024	0.038	0.821	335
Qwen7B	0.182	0.118	0.143	0.774	335
All	0.201	0.327	0.215	0.635	335
WT					
LLaMa70B	0.241	0.433	0.310	0.310	336
LLaMa8B	0.328	0.842	0.472	0.074	336
ModBert	0.407	0.717	0.520	0.613	258
OLMo32B	0.292	0.466	0.359	0.465	331
OLMo7B	0.348	0.958	0.511	0.009	336
Qwen32B	0.267	0.258	0.263	0.601	336
Qwen7B	0.264	0.392	0.315	0.455	336
All	0.307	0.581	0.393	0.361	336
Fora					
LLaMa70B	0.333	0.058	0.099	0.760	382
LLaMa8B	0.386	0.584	0.465	0.562	382
ModBert	0.409	0.631	0.497	0.599	2764
OLMo32B	0.391	0.066	0.113	0.765	382
OLMo7B	0.358	0.993	0.526	0.008	382
Qwen32B	0.438	0.051	0.092	0.773	382
Qwen7B	0.412	0.102	0.164	0.759	382
All	0.390	0.355	0.279	0.604	382
IQ2					
LLaMa70B	0.265	0.137	0.181	0.703	374
LLaMa8B	0.318	0.695	0.436	0.290	374
ModBert	0.399	0.634	0.490	0.572	2393
OLMo32B	0.273	0.092	0.137	0.736	374
OLMo7B	0.345	0.977	0.510	0.000	374
Qwen32B	0.167	0.038	0.062	0.743	374
Qwen7B	0.369	0.290	0.325	0.693	374
All	0.305	0.409	0.306	0.534	374
WHoW					
LLaMa70B	0.312	0.159	0.211	0.721	364
LLaMa8B	0.349	0.778	0.482	0.343	364
ModBert	0.418	0.626	0.502	0.597	1816
OLMo32B	0.455	0.079	0.135	0.779	364
OLMo7B	0.349	1.000	0.517	0.025	364
Qwen32B	0.296	0.063	0.105	0.762	364
Qwen7B	0.370	0.294	0.327	0.697	364
All	0.364	0.428	0.326	0.561	364
All datasets					
LLaMa70B	0.253	0.197	0.194	0.659	2000
LLaMa8B	0.313	0.677	0.425	0.388	2000
OLMo32B	0.285	0.149	0.162	0.716	1994
ModBert	0.400	0.555	0.478	0.609	2000
OLMo7B	0.316	0.986	0.476	0.007	2000
Qwen32B	0.236	0.105	0.121	0.728	2000
Qwen7B	0.296	0.229	0.242	0.686	2000
All	0.300	0.485	0.300	0.542	2000

Table 4: Classifier performance for facilitator prediction using positive-class precision, recall and F1 ($F1_p$) and negative-class F1 ($F1_n$). Rows are struck through when either $F1_p$ or $F1_n$ collapses (< 0.15). Best non-collapsed $F1_p$ and $F1_n$ values for each dataset split are highlighted in bold. “All” rows report macro averages across models.

arate methodologies (see §2). We frame the task as a classification problem, aiming to identify the last comment made by a non-facilitator before a facilitator intervenes. To this end, we evaluate six open-source LLMs (App. C.1) and train a Modern-

Bert classifier (Warner et al., 2025).¹⁰ Participant names are redacted, and each instance includes the three preceding comments as context. We use a 60–20–20 split for training the classifier, and choose 2,000 randomly sampled instances from the same test set for the LLM evaluation, evenly distributed across the sub-datasets of PEFK. We use a single Quadro RTX 6000 GPU for the classifier and two GPUs for the LLMs, collectively requiring about 72 hours of computation; see further details in App. C.2.

Deciding when NOT to intervene is easy. Table 4 shows the results for the facilitator intervention prediction task. We find that while most models are adept at identifying cases where facilitators did not intervene (large $F1_n$), they are generally much less successful at identifying when they *would* (poor $F1_p$) mirroring the fact that humans are more certain of when not to intervene (§4.2).

Finding 3. *Deciding when not to facilitate is easy both for humans and LLMs. Deciding when to facilitate, is not.*

Do LLMs over or under-intervene? LLM behavior in facilitation is polarized: some models exhibit a pervasive tendency to constantly facilitate (reflected by extremely low $F1_n$), a pattern documented in prior work (Tsirmpas et al., 2026). Conversely, we observe that other models rarely intervene (evidenced by extremely low $F1_p$ —both crossed out in Table 4), a behavior that has been largely overlooked in the existing literature. We also note that discussion domain systematically influences LLMs but not humans in this regard—see App. F.5.

Are EncOn models the solution? We find that the ModernBert classifiers can generally compete with LLMs in this task, and seem much more consistent, outperforming them in three out of the five datasets, and scoring the highest $F1_p$ scores across all datasets.¹¹ We posit that overall, these classifiers may be a much more scalable, reliable and viable solution compared to LLMs for the task, even though they have been overlooked in current

literature.¹²

Finding 4. *Out-of-the-box LLMs are generally not consistent in deciding when to facilitate, in contrast to traditional EncOn classifiers.*

6 Conclusion

We investigated how the timing of humans and LLMs differs in facilitative interventions. By conducting a literature review, we isolated three main definitions of facilitative interventions. We then conducted an survey task using facilitative social science experts on a standardized, aggregated corpus of all relevant facilitation datasets, finding that human facilitators tend to adopt a cautious approach, are more confident when deciding not to intervene than when deciding to intervene, and may disagree because they focus on different spans of the discussion. We are the first to conduct extensive experiments on facilitative intervention prediction using LLMs, finding that they tend to over-intervene, skew towards positive reinforcement when the situation is ambiguous, and remain inconsistent across facilitation decisions, while simpler ModernBERT classifiers trained on existing datasets can outperform them.

Despite optimism that LLMs would revolutionize facilitation, our findings suggest that progress may be incremental, not transformative—LLMs can not meaningfully facilitate until we move beyond limited datasets. Previous approaches—such as focusing on escalation or general moderation—provided large, scalable data, but only captured surface-level behaviors. They failed to give us a true understanding of how human facilitators operate, a complex and open question that can not be answered without annotated data that captures genuine facilitative dynamics.

Limitations

This work is not without limitations. With additional resources, we could have incorporated a larger survey or, given that we understood the task’s subjectivity more adequately, even annotation schemata (e.g., sentence-by-sentence, apart from the excerpt) and explored a wider range of prompting settings. Also the opposite is true; a possible criticism of our decision is that we do not pro-

¹⁰ Due to substantial differences in text format, we train separate models for oral and written datasets (App. B.3); an ablation study using a single model across all datasets shows reduced performance (App. F.3).

¹¹ There technically exists ways of boosting the performance of LLM classifiers, but this is outside the scope of this work—see App. C.1.

¹² Another notable advantage of these models is that we can tune the precision-recall tradeoff by adjusting the probability threshold over which the model determines that an intervention should happen, thus configuring how proactive or cautious they are—see App. F.4.

vide the full discussion thread, which in some cases could be extremely large and impractical to include in its entirety. Therefore, this organic limitation led to us to choose a middle ground, that is providing excerpts of the discussions. Facilitation decisions in real-world settings are often shaped by long-term interaction dynamics, participant history, community norms, and evolving conversational context, which may not be fully captured in the excerpts.

Furthermore, because providing rationales was optional for participants, we collected very few, resulting in the loss of a potentially valuable resource for interpretation—this information is valuable and should be incorporated in future tasks concerning facilitation. We only realized their usefulness during the analysis of our results, although it has to be acknowledged that making them mandatory would have significantly added to the workload of the participants.

Moreover, the intervention strategies examined here are only two: positive and negative reinforcement. Ideally, the task should be conducted without the use of different intervention labels or very fine-grained labels, as prior work has shown that different domains call for different facilitative tactic taxonomies (Park et al., 2012; Chen et al., 2025; Schroeder et al., 2024). The decision of including the distinction between negative and positive reinforcement was motivated by the fact that the participants were over-focusing on the moderation part during the refinement stage of the task. We also acknowledge that different discussion domains may favor different facilitative norms, which may have impacted the performance of our LLM facilitative classifiers. However, these domain-specific norms have not been explored in the literature, and the elicitation of optimal prompts is not our primary research objective, and thus considered outside of the scope of this work.

Additionally, although early work attempted to estimate interventions in terms of external metrics such as toxicity (Hua et al., 2018; Chang and Danescu-Niculescu-Mizil, 2019; Yuan and Singh, 2023), recent work has shown that facilitators may have different objectives depending on the form of discussion, which may require the use of different evaluation metrics for each domain (Korre et al., 2025). Since our work attempts to generalize interventions across multiple domains, we can not rely on external signals, but only on whether an intervention would (or did) occur.

Finally, the datasets included in PEFK are pre-

dominantly English-language and drawn from specific online or institutional discussion settings. As a result, the findings may not generalize to multilingual environments and culturally distinct moderation norms.

Ethical Considerations

All experiments and data have passed through official approval and were overseen by an Ethics Committee. The participants were selected via formal recruitment methods based on familiarity with the task and field, and were compensated fairly with a monthly salary of 1000€. ¹³ They signed informed consent forms before beginning the survey, were informed and consented to being exposed to potentially offensive or inappropriate content, and were allowed to skip uncomfortable tasks at any point, with no required justification. All necessary data were anonymized both during data collection and in the final, presented artifacts, while non-essential information was destroyed with the conclusion of the survey in accordance with GDPR (European Union, 2016) and the EU AI Act (European Union, 2024). Participant demographic data was not collected, since it was deemed unnecessary for the purposes of this study.

In order to help research in the field of facilitation, we openly release the code that constructs PEFK under a GPLv3 license. All datasets used in PEFK are obtained from the sources provided by the respective authors and are consistent with intended use (academic research). We do not distribute the corpus separately, and thus no additional data license is applicable. The licenses of the original distribution apply. The Fora dataset (Schroeder et al., 2024) is not publicly available because it may contain personally identifying information; researchers should contact the authors of Fora for access. Finally, we use exclusively open-source models, libraries and LLMs, in order to ensure reproducibility, protect user data, and comply with relevant regulations (see App. C.1).

We used LLMs for some orthographic and grammatical support throughout the document, as well as for code generation in some instances. All outputs have been verified by the authors.

¹³Specific information on posted requirements and the recruitment process will be made available after review, since they would compromise anonymity.

References

- Fazl Barez, Tung-Yu Wu, Iván Arcuschin, Michael Lan, Vincent Wang, Noah Siegel, Nicolas Collignon, Clement Neo, Isabelle Lee, Alasdair Paren, and 1 others. 2025. Chain-of-thought is not explainability. *Preprint, alphaXiv*, page v1.
- James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. 2024. Synthetic replacements for human survey data? the perils of large language models. *Political Analysis*, 32(4):401–416.
- Jason W. Burton, Ezequiel Lopez-Lopez, Shahar Hechtlinger, Zoe Rahwan, Samuel Aeschbach, Michiel A. Bakker, Joshua A. Becker, Aleks Berditchevskaia, Julian Berger, Levin Brinkmann, Lucie Flek, Stefan M. Herzog, Saffron Huang, Sayash Kapoor, Arvind Narayanan, Anne-Marie Nussberger, Taha Yasseri, Pietro Nickl, Abdullah Almaatouq, and 9 others. 2024. How large language models can reshape collective intelligence. *Nature Human Behaviour*, 8(9):1643–1655.
- Jonathan P. Chang and Cristian Danescu-Niculescu-Mizil. 2019. Trouble on the horizon: Forecasting the derailment of online conversations as they develop. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4743–4754, Hong Kong, China. Association for Computational Linguistics.
- Ming-Bin Chen, Lea Frermann, and Jey Han Lau. 2025. WHOw: A cross-domain approach for analysing conversation moderation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 2091–2126, Albuquerque, New Mexico. Association for Computational Linguistics.
- Andy Cockburn, Philip Quinn, and Carl Gutwin. 2017. The effects of interaction sequencing on user experience and preference. *International Journal of Human-Computer Studies*, 108:89–104.
- William D Crano. 1977. Primacy versus recency in retention of information and opinion change. *The Journal of Social Psychology*, 101(1):87–96.
- Stefano Cresci, Amaury Trujillo, and Tiziano Fagni. 2022. Personalized interventions for online moderation. In *Proceedings of the 33rd ACM Conference on Hypertext and Social Media*, HT ’22, page 248–251, New York, NY, USA. Association for Computing Machinery.
- Christine De Kock, Tom Stafford, and Andreas Vlachos. 2022. How to disagree well: Investigating the dispute tactics used on wikipedia. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3824–3837, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Christine De Kock and Andreas Vlachos. 2021. I beg to differ: A study of constructive disagreement in online conversations. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2017–2027, Online. Association for Computational Linguistics.
- European Union. 2016. Regulation (EU) 2016/679 of the European Parliament and of the Council.
- European Union. 2024. Regulation (eu) 2024/1689 of the european parliament and of the council of 13 june 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>. Official Journal of the European Union, L 2024/1689.
- Neele Falk, Iman Jundi, Eva Maria Vecchi, and Gabriella Lapesa. 2021. Predicting moderation of deliberative arguments: Is argument quality the key? In *Proceedings of the 8th Workshop on Argument Mining*, pages 133–141, Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Neele Falk, Eva Vecchi, Iman Jundi, and Gabriella Lapesa. 2024. Moderation in the wild: Investigating user-driven moderation in online discussions. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 992–1013, St. Julian’s, Malta. Association for Computational Linguistics.
- Alvan R. Feinstein and Domenic V. Cicchetti. 1990. High agreement but low kappa: I. the problems of two paradoxes. *Journal of Clinical Epidemiology*, 43:543–549.
- Rena Gao, Ming-Bin Chen, Lea Frermann, and Jey Han Lau. 2025. Moderation matters: Measuring conversational moderation impact in English as a second language group discussion. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 2070–2095, Vienna, Austria. Association for Computational Linguistics.
- Jim Hewitt. 2003. How habitual online practices affect the development of asynchronous discussion threads. *Journal of Educational computing research*, 28(1):31–45.
- Yiqing Hua, Cristian Danescu-Niculescu-Mizil, Dario Taraborelli, Nithum Thain, Jeffery Sorensen, and Lucas Dixon. 2018. Wikiconv: A corpus of the complete conversational history of a large online collaborative community. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2818–2823, Brussels, Belgium. Association for Computational Linguistics.
- Katerina Korre, Dimitris Tsirmpas, Nikos Gkoumas, Emma Cabalé, Danai Myrtzani, Theodoros Evgeniou, Ion Androutsopoulos, and John Pavlopoulos.

794	2025. Evaluation and facilitation of online discussions in the LLM era: A survey . In <i>Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing</i> , pages 24443–24462, Suzhou, China. Association for Computational Linguistics.	847
795		848
796		849
797		850
798		851
799		852
800	Vinay Koshy, Frederick Choi, Yi-Shyuan Chiang, Hari Sundaram, Eshwar Chandrasekharan, and Karrie Karahalios. 2024. Venire: A machine learning-guided panel review system for community content moderation . <i>CoRR</i> , abs/2410.23448.	853
801		854
802		855
803		
804		
805	Anna Sophie Kümpel and Julian Unkel. 2021. (why) does comment presentation order matter for the effects of user comments? assessing the role of the availability heuristic and the bandwagon heuristic. <i>Communication Research Reports</i> , 38(4):217–228.	856
806		857
807		858
808		859
809		860
810	Cliff Lampe, Paul Zube, Jusil Lee, Chul Hyun Park, and Erik Johnston. 2014. Crowdsourcing civility: A natural experiment examining the effects of distributed moderation in online forums. <i>Government Information Quarterly</i> , 31(2):317–326.	861
811		
812		
813		
814		
815	Jorge Nathan Matias. 2019. The civic labor of volunteer moderators online . <i>Social Media + Society</i> , 5.	862
816		863
817	Deokgun Park, Simranjit Sachar, Nicholas Diakopoulos, and Niklas Elmqvist. 2016. Supporting comment moderators in identifying high quality online news comments. In <i>Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems</i> , pages 1114–1125.	864
818		865
819		
820		
821		
822		
823	Joonsuk Park, Sally Klingel, Claire Cardie, Mary Newhart, Cynthia Farina, and Joan-Josep Vallbé. 2012. Facilitative moderation for online participation in erulemaking . In <i>Proceedings of the 13th Annual International Conference on Digital Government Research</i> , dg.o '12, page 173–182, New York, NY, USA. Association for Computing Machinery.	866
824		867
825		868
826		869
827		870
828		871
829		
830	Perspective API. 2026. Perspective api . Accessed: 2026-02-24.	872
831		873
832	Luca Rossi, Katherine Harrison, and Irina Shklovski. 2024. The problems of llm-generated data in social science research . <i>Sociologica</i> , 18(2):145–168.	874
833		875
834		876
835	Paul Röttger, Bertie Vidgen, Dirk Hovy, and Janet Pierrehumbert. 2022. Two contrasting data annotation paradigms for subjective NLP tasks . In <i>Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 175–190, Seattle, United States. Association for Computational Linguistics.	877
836		878
837		879
838		880
839		881
840		882
841		883
842		884
843	Manish Sanwal. 2025. Layered chain-of-thought prompting for multi-agent llm systems: A comprehensive approach to explainable large language models. <i>arXiv preprint arXiv:2501.18645</i> .	885
844		886
845		887
846		888
		889
	Brennan Schaffner, Arjun Nitin Bhagoji, Siyuan Cheng, Jacqueline Mei, Jay L Shen, Grace Wang, Marshini Chetty, Nick Feamster, Genevieve Lakier, and Chenhao Tan. 2024. "community guidelines make this the best party on the internet": An in-depth study of online platforms' content moderation policies . In <i>Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems</i> , CHI '24, New York, NY, USA. Association for Computing Machinery.	890
		891
		892
		893
		894
		895
		896
		897
		898
		899
		900
	Hope Schroeder, Deb Roy, and Jad Kabbara. 2024. Fora: A corpus and framework for the study of facilitated dialogue . In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 13985–14001, Bangkok, Thailand. Association for Computational Linguistics.	901
		902
		903
		904
	Joseph Seering. 2020. Reconsidering self-moderation: the role of research in supporting community-based models for online content moderation . <i>Proc. ACM Hum.-Comput. Interact.</i> , 4(CSCW2).	
	Tsukasa Shiota, Takashi Yamamura, and Kazutaka Shimada. 2018. Analysis of facilitators' behaviors in multi-party conversations for constructing a digital facilitator system. In <i>Collaboration Technologies and Social Computing</i> , pages 145–158, Cham. Springer International Publishing.	
	Christopher T. Small, Ivan Vendrov, Esin Durmus, Hadjar Homaei, Elizabeth Barry, Julien Cornebise, Ted Suzman, Deep Ganguli, and Colin Megill. 2023. Opportunities and risks of llms for scalable deliberation with polis . <i>ArXiv</i> , abs/2306.11932.	
	Amaury Trujillo and Stefano Cresci. 2022. Make reddit great again: Assessing community effects of moderation interventions on r/the_donald . <i>Proc. ACM Hum.-Comput. Interact.</i> , 6(CSCW2).	
	Dimitris Tsirmpas, Ion Androustopoulos, and John Pavlopoulos. 2026. Designing synthetic discussion generation systems: A case study for online facilitation . <i>Preprint</i> , arXiv:2503.16505.	
	Miles Turpin, Julian Michael, Ethan Perez, and Samuel Bowman. 2023. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. <i>Advances in Neural Information Processing Systems</i> , 36:74952–74965.	
	Benjamin Warner, Antoine Chaffin, Benjamin Clavié, Orion Weller, Oskar Hallström, Said Taghadouini, Alexis Gallagher, Raja Biswas, Faisal Ladhak, Tom Aarsen, Griffin Thomas Adams, Jeremy Howard, and Iacopo Poli. 2025. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 2526–2547, Vienna, Austria. Association for Computational Linguistics.	
	Jiaqing Yuan and Munindar P. Singh. 2023. Conversation modeling to predict derailment . <i>Proceedings of the International AAAI Conference on Web and Social Media</i> , 17(1):926–935.	

Justine Zhang, Ravi Kumar, Sujith Ravi, and Cristian Danescu-Niculescu-Mizil. 2016. [Conversational flow in Oxford-style debates](#). In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 136–141, San Diego, California. Association for Computational Linguistics.

A Acronyms

NLP	Natural Language Processing
NLU	Natural Language Understanding
ML	Machine Learning
RQ	Research Question
PEFK	Prosocial and Effective Facilitation in Konversations
EncOn	Encoder-Only
IAA	Inter-Annotator Agreement
UMOD	User Moderation
IQ2	Intelligence Squared
CeRI	Cornell eRulemaking Initiative
WHoW	Why-How-Who
WT	WikiTactics
CoT	Chain of Thought

B Dataset details

B.1 Preprocessing

General

- We exclude comments with no text.
- We exclude discussions with less than two distinct participants.
- We exclude discussions which are common between wiktactics and wikiconv as well as wikidisputes and wikiconv.
 - There may be duplicate discussions between wikidisputes and wiktactics, but we allow them since they feature complementary information.

Wikiconv

- The Wikiconv corpus does not contain information about which user is a moderator/facilitator. Therefore, all comments relating to Wikiconv are tagged as non-moderators.
- Additionally, we follow the instructions of the original researchers ([Hua et al., 2018](#)), and select only discussions which have at least two comments by different users.
 - Wikipedia (thankfully) does not track users who log in with only an IP address (in the original dataset, their `user_id` is always set to 0 and their username is of the form 211.111.111.xxx. We consider each such username to be a separate user.
 - Due to the size of the dataset, we have to partially load it during preprocessing. Thus, there is a small chance every 100,000 records that a discussion is marked as a false negative and a part of it gets discarded.
 - We only include English comments in the final dataset. We use a small, efficient library (`py3langid`) for language recognition, due to the large size of Wikiconv. Non-English comments are discarded before selecting valid discussions (see point above).

Wiktactics

- We infer facilitative actions by whether the comment belongs in any of the following categories:
 - Asking questions
 - Coordinating edits
 - Providing clarification
 - Suggesting a compromise
 - Contextualisation
- The above tactics are a subset of the Coordinative labels used in the WikiTactics paper. They were selected because they are not used necessarily on 1-1 discussions; they could reasonably be applied by third-party participants. Contrast them with other Coordinative labels such as “Conceding/recanting” and “I don’t know”.

Wikidisputes

- Since only 0.03% of the comments in the dataset are made by moderators, we mark the dataset as not supporting moderation.

Name	Type	Description
conv_id	string	The discussion’s ID. Comments under the same discussion refer to the same discussion ID.
message_id	string	The message’s (comment’s) unique ID.
reply_to	string	The ID of the comment which the current comment responds to. nan if the comment does not respond to another comment (e.g., it’s the Original Post (OP)).
user	string	Username or hash of the user that posted the comment
is_moderator	bool	Whether the user is a facilitator.
moderation_supported	bool	True if the moderation labels are directly computed from the original dataset
escalated	bool	A discussion-level measure denoting discussions which have been derailed
escalation_supported	bool	True if the escalation labels are directly computed from the original dataset
text	string	The contents of the comment
dataset	string	The dataset from which this comment originated from
notes	JSON	A dictionary holding notable dataset-specific information
toxicity	float	The “toxicity” score given to the comment by the Perspective API
severe_toxicity	float	The “severe toxicity” score given to the comment by the Perspective API

Table 5: The full set of columns used for each comment in the [PEFK](#) dataset.

UMOD

- Facilitative actions are marked as a gradient from 0 (no facilitation) to 1 (full facilitation). We adopt a threshold of 0.75 to consider an action as facilitative, with more than 50% participant agreement (measured as entropy in the original dataset).

CMV-AWRY2

- We mark a discussion as escalated when the derailment value (from the official dataset) is in the 60th upper percentile.
- We remove deleted (“[deleted]”) comments.

B.2 Schema

Table 5 details the information available for each comment in the dataset. Some information, such as whether the person talking is a moderator / facilitator, is available only on some sub-datasets. Also note that the `is_moderator` column may hold different definitions of what we consider as moderator / facilitator. UMOD for instance considers facilitative comments made by users, while others such as Fora consider a comment facilitative if it was made by a dedicated facilitator.




Figure 5: Percentage of comments made by professional facilitators (see Def. 2) for each of the sub-datasets of [PEFK](#). Bars denote 95% confidence intervals. Despite the different kinds of facilitators in online platforms and various real-life-discussion domains, the overall rate of facilitation remains similar across datasets.

B.3 Exploratory Data Analysis

We explore some initial frequencies with regard to facilitator comments in [PEFK](#), comparing its oral vs. written subsets, as well as vocabulary discrepancies.

Oral vs. written datasets in PEFK There currently exists a rift in facilitation research; while some previous work concerns itself with forums and online discussions ([Park et al., 2012](#)), others are exclusively focused on live, real-life, oral discussions ([Schroeder et al., 2024](#); [Chen et al., 2025](#)).




Figure 6: **Left:** Comments per discussion on the PEFK dataset. Note the significant increase in comments made in transcribed datasets such as Fora and IQ2. This deviation is explained by the fact that a single human utterance is regularly broken up to many comments in order to simulate interruptions or participants talking over each other. **Right:** Number of words per comment in all discussions in PEFK, grouped by original dataset. Comments across written datasets follow the same distribution, while oral datasets feature a significant number of comments with very few words.

This rift is an issue for facilitation research, since oral and online written discussions feature fundamentally different kinds of facilitation interventions, as shown by stark differences in proposed facilitation schemas between online forums (Park et al., 2012), and real-life discussions (Schroeder et al., 2024; Chen et al., 2025). Additionally, the data themselves are different; posts and comments are self-contained, while transcripts regularly feature interruptions, participants talking in parallel (see App. B.3), and linguistic features that are necessary for coordination (e.g., “Thank you.”, “I agree.”). It is worth noting that in our case, all written discussions are asynchronous (the participants may join and leave at any time), while all oral discussions are synchronous (the participants are all active at the same time).

Figure 6 shows some of the structural differences of oral and written datasets. oral datasets feature a significantly larger number of comments per discussion, since instances where the participants talk over each other, or interrupt the other are transcribed as separate comments. Additionally, oral datasets feature a large amount of comments with very few words, since a lot of utterances are short but necessary for social co-ordination (e.g., “Thank you”, “Yes”, “I agree”). Figure 5 shows the percentage of comments made by facilitators for each of our datasets. In almost all datasets, a third of the comments are made by facilitators.¹⁴

¹⁴We do not include the UMOD dataset, since there is no established ground-truth (see App. B.1).

Vocabulary discrepancies in PEFK datasets.

We also the data to include only messages written by moderators, then cleaned the text by removing URLs, punctuation, and other noise. After preprocessing, we used a TF-IDF vectorization approach with n-grams (bigrams and trigrams) to extract phrases. Crucially, the TF-IDF model was trained on all datasets combined, so that common, generic phrases were downweighted while phrases that are more distinctive to each dataset were emphasized. We then computed the average TF-IDF score of each phrase within each dataset and selected the top-ranking ones as representative expressions of moderator language in that context. The results can be found in Figure 7.

The most important phrases in CeRI suggest a formal, policy-oriented environment with an emphasis on structured participation and regulatory discussion (e.g., references to “proposed rule” and “welcome regulations”). In *fora*, the language is more conversational and facilitative, with phrases such as “great thank”, “feel free”, and “awesome thank” indicating a more informal moderation style. The IQ2 dataset reflects the debate setting, with phrases like “ladies gentlemen” and “arguing motion”. In UMOD, moderators engage more in argumentation and clarification, as seen in phrases like “change view”, “missing point” reminiscing more facilitative settings and suggesting a focus on reasoning and discussion quality. The WHoW dataset appears similar to moderated discussions or panels, with phrases like “question sir” and “answer question” indicating turn-taking and audience interaction. Finally, WT shows a highly procedural




Figure 7: Most common moderator bigrams in PEFK per dataset. Function words are excluded.

Table 6 shows details on the LLMs used for this study. We specifically use 4-bit quantized versions from the unsloth huggingface repository.

Family	Version	#Params
OLMo	2 (0325)	32B
Qwen	2.5	32B
LLaMa	3.3	70B
OLMo	3	7B
Qwen	2.5	7B
LLaMa	3.1	8B

Table 6: LLMs used in this study.

and rule-based style, with phrases such as “reliable sources”, “original research” and “talk page” reflecting the norms and policies of collaborative editing environments like Wikipedia.

C Models

C.1 LLMs

We use only open-source models, as proprietary models such as OpenAI’s GPT-5 yield much less reproducible results (Bisbee et al., 2024). Furthermore, the use of open-source models requires much fewer resources, enabling research on this task with a less demanding budget, as well as scalable deployment in real-world settings (Tsirmpas et al., 2026). It is also worth noting that since a part of our dataset (Fora) is not licensed, we can not risk leaking the data contained within to proprietary models. Generally, the use of local models protects the personal data of the users contained in our datasets, ensures that our research can be replicated in domains where anonymity must be maintained, and renders our work compatible with existing legislation (European Union, 2024, 2016).

Proprietary LLMs such as the OpenAI GPT family may perform better, but cost several orders of magnitude more, and are thus not a realistic solution (Tsirmpas et al., 2026). While the performance of LLM classifiers could be reasonably increased by providing few-shot examples or finetuning, these approaches would significantly bias the model since no concrete definition of facilitation exists (Korre et al., 2025), and the examples would have to be different for each different discussion domain in PEFK—which goes against our motivation of discovering patterns between oral and written datasets.

C.2 Transformers

We construct our training and evaluation data from a unified corpus of threaded discussions, where each instance corresponds to a target comment augmented with its conversational context. For every target comment, we retrieve up to $K = 3$ preceding comments along the reply chain. Each comment is truncated to a maximum of 3000 characters to control sequence length. The input is formatted using XML-style tags, where context comments are wrapped with <CTX> and the target comment with <TGT>. The final sequence is formed by concatenating context comments in chronological order followed by the target comment. We consider binary classification tasks (e.g., is_moderator, should_intervene), with labels cast to floating point values.

We fine-tune a pretrained transformer encoder (ModernModBert-large) with an added classification head that outputs a single logit. The task is formulated as a single-label instance of multi-label classification, enabling the use of a sigmoid activation.

Training is conducted using the Hugging Face Trainer framework with several modifications. We optimize the model using binary cross-entropy with logits:

$$\mathcal{L} = \text{BCEWithLogitsLoss}(\mathbf{z}, \mathbf{y}; \text{pos_weight}),$$

where $\mathbf{z}$ are the predicted logits and $\mathbf{y}$ are the ground-truth labels. To account for class imbalance, the positive class is weighted as. We freeze the base transformer parameters and train only the classification head in order to reduce computational cost and limit overfitting.

To improve computational efficiency, we employ a bucketed batching strategy that groups instances of similar length (approximated in characters) into batches of size 64. This reduces padding and improves throughput. Batches are shuffled at the bucket level at each epoch.

Training is performed for up to 80 epochs, with evaluation, logging, and checkpointing conducted once per epoch. The best model is selected based on validation loss. We apply early stopping with a warm-up phase: monitoring begins after one evaluation step, and training stops if validation loss does not improve by at least $\delta = 0.001$ for 6 consecutive evaluation steps. Model performance is evaluated using accuracy and binary F1 scores.

All experiments are implemented in PyTorch using the Hugging Face Transformers library. Data loading uses 4 worker processes, and GPU acceleration is utilized when available. We fix the random seed to 42 to ensure reproducibility.

D Human survey procedure

D.1 Facilitative training

Before the main survey on facilitative intervention, all participants were extensively trained on currently established facilitative tactics. This task aimed to familiarize the participants with what constitutes a facilitative intervention, as well as how professional facilitators operate in real-life scenarios.

Specifically, we split the participants into three groups, each of which was assigned a different taxonomy of facilitative interventions (Chen et al., 2025; Schroeder et al., 2024; Park et al., 2012). The participants were tasked with classifying 830 facilitated comments from a sample of continuous discussions in PEFK. The guidelines and results of this annotation are not shown in this study, but they are available upon request.

D.2 Positive vs. negative reinforcement

The distinction between positive/negative reinforcement is present to suggested facilitative tactics taxonomies (Table 2), including prior work such as Fora (Schroeder et al., 2024) WHoW (Chen et al., 2025), and CeRI (Park et al., 2012). We conservatively estimate whether the tactics described in each taxonomy correspond to our definitions of positive and negative reinforcement.

The facilitative taxonomy proposed in Fora does not include any tactics that could be unambiguously

Taxonomy & Tactic	Type
Fora	
Express appreciation	+
Express agreement or affirmation	+
Open invitation to participants	
Specific invitation to participant	
Model examples	
Follow-up question	
Make connections	
WHoW	
Informational	
Coordinative	-
Social	+
Probing	
Confronting	
Instruction	-
Interpretation	
Supplement	
CeRI	
Welcoming	+
Encouragement / appreciation	+
Maintaining civil discourse	-
Explaining why comment not removed	
Indicating irrelevant/off-point comments	-
Pointing to relevant information	
Pointing out effective commenting	+
Asking for more information/sources	
Asking for solutions/alternatives	
Encourage engagement with others	
Pose question to community	

Table 7: Facilitation taxonomies (bold) with associated tactics (normal text) and reinforcement type. +: positive reinforcement, -: negative reinforcement, empty if the tactic does not clearly correspond to any category.

considered as negative reinforcement, which is expected, since facilitators deployed in the relevant study were supposed to only act as coordinators. Notably, several tactics may simultaneously exhibit both positive and negative characteristics across all taxonomies, depending on context and the exact implementation used by the facilitator. We therefore use these taxonomies only as guidance during annotation, and refer the reader to the original works for a more comprehensive description and analysis of facilitation strategies.

D.3 Label Transformation

To efficiently calculate intra-participant consistency, we transform the fine-grained, ordinal label per type adopting a multiclass schema. We define a threshold of ≥ 3 and then using Table 8 we relabel the items.

Positive	Negative	Neither	Final
✗	✗	✗	None (0)
✓	✗	✗	Positive (1)
✗	✓	✗	Negative (2)
✗	✗	✓	None (0)
✓	✓	✗	Both (3)
✓	✗	✓	Positive (1)
✗	✓	✓	Negative (2)
✓	✓	✓	Both (3)

Table 8: Reinforcement label transformations.




Figure 8: Label distributions of human participants per dataset.

D.4 Are facilitative tendencies dataset-dependent?

Facilitative tendencies generally follow a conservative pattern, with no reinforcement selected most often. Despite the differences between the datasets (forums, debates, and radio podcasts), we observe similar trends in the choice between positive and negative reinforcement, with positive reinforcement being more common overall. Two exceptions are *cmv_awry* and *wikiconv*, where negative interventions are selected more often than positive ones, by 15% and 1%, respectively (Fig. 8).

D.5 Participant consistency

To evaluate the consistency of individual participants, we measured intra-participant agreement using purposefully repeated instances of identical discussions.

Agreement was computed as a strict match in the labeling of the same instance (duplicated pair). For each duplicated pair, agreement was assigned a value of 1 if both labels were identical and 0 otherwise in order to calculate a simple pairwise agreement. Hence, the final intra-participant agree-

participant	Agreement (%)
A2	100.00
A3	66.67
A4	88.89
A5	100.00
A6	100.00
A7	100.00
A8	100.00
A9	100.00
A10	100.00
A11	100.00

Table 9: Intra-participant percentage agreement (consistency) over all subdatasets of PEFK.

ment for each participant was then calculated as the mean of these binary agreement scores. Since our classes are heavily skewed and imbalanced, traditional agreement metrics tend to collapse, showing extremely low scores even for almost-100% agreement (Röttger et al., 2022; Feinstein and Cicchetti, 1990). For the purposes of simplicity, we thus run our analysis in this part with simple percentage agreement.

We report results in Table 9. Most participants demonstrated near-perfect consistency, with several achieving 100% agreement. Participant A3 is the only participant who showed reduced consistency (66.67%), suggesting potential variability in interpretation or labeling behavior.

E Participants Selection and Coordination

Selection Participant selection required a response to the public call, followed by a round of interviews to verify that candidates fulfilled the criteria for the role. All successful participants held a BA in Languages, Communication, or History, and an MSc in either Digital Humanities or Language Technology. They also had experience with NLP annotation tasks, such as sentiment analysis and toxicity detection. We would also like to highlight that controls were conducted before the survey (hiring), during the survey (training and discussion), and after the survey (removal of outlier or inconsistent participants), to ensure the highest quality of results.

Coordination The coordination of the project was largely based on weekly meetings that allowed the integration of participant feedback and potential

Ph.	Dates	Activity
1	Dec	Recruitment
2	Dec–Jan	Training on facilitation strategies
3	Jan 1–9	Onboarding
4	Jan 9–15	Pilot
5	Mid Jan–Mid Feb	Main experiment and discussion
6	End Feb–Mid Mar	Main experiment completion
1–6	Weekly	Discussion & feedback rounds

Table 10: Coordination timeline.

changes to our survey methodology. An overview of the project timeline can be found in Table 10. Before the commencement of the main survey, two crucial steps were carried out. The first was familiarizing participants with facilitation tactics; more information can be found in Section D.1. The second was a pilot round with ~ 160 instances, which allowed us to better understand the challenges of the task. For example, one major concern was time. Depending on instance length, labeling took between 2 and 5 minutes. This led us to break the task into smaller subtasks, spread across almost 3 months (see Phases 5 and 6 in Table 10), with annotators labeling about 400 instances per round.

F Additional Experiments

F.1 Can LLMs understand what facilitation is?

One way of establishing whether the models can understand the concept of facilitation is exploring whether these models can distinguish between facilitative and non-facilitative comments in the first place. Their performance in this task should give us an indication of how well-defined the task is in relevant NLP datasets. We thus explore “*facilitator detection*” a Natural Language Understanding (NLU) task which aims to identify whether a comment was made by a facilitator (Def. 2). This task has been explored exactly once by Shiota et al. (2018), although the authors used hand-crafted features instead of relying exclusively on the language used by the facilitators, and used only role-play focused scenarios instead of real-life facilitation ones.

Table 14 presents the performance of these same models in the task of facilitation detection. While we observe instances of strong performance—such as Qwen32B achieving high scores on CeRI ($F1_p=0.904$, $F1_n=0.971$)—the models generally do not score reliably across diverse datasets (e.g., LLaMa70B in WT yields $F1_p=0.500$). This

suggests that the fundamental cause of under-performance is less a confusion with the traditional moderator role and more a structural issue with the datasets themselves. Prior work utilizing traditional ML (Shiota et al., 2018), LLMs, and EncOn classifiers shows similar challenges on other tasks (Chen et al., 2025), indicating that this finding is not merely a result of our specific labeling scheme or dataset selection.¹⁵ Thus, significant, generalizable improvements can only be realized by incorporating larger professional facilitation datasets, where interventions can be *reliably defined and distinguished* from non-facilitative comments, or by developing larger-scale corpora that contain annotated *facilitative comments* (Def. 3)—an extremely difficult task given the inherent subjectivity of the task (§4).

F.2 LLM sensitivity ablations

We run an ablation study using five instruction prompt variants on a 20% sample of the dataset. We generated these variants via a local Gemma-4 model (specifically, the <https://huggingface.co/unsloth/gemma-4-E4B-it-GGUF> variant, generated using the TextGen open-source software—<https://github.com/oobabooga/textgen>) in order to avoid researcher bias, and verified the resulting output. Two of these variants include specific instructions warning against over-intervention (“cautious 1,2”). The pairwise agreement scores for four models can be found in Table 11 and the instruction variants in App. H.2. We find that LLM judgments seem to remain broadly consistent, even with large changes to the LLM input for the specific task. Some inconsistencies do exist, as is the case with all LLM-related tasks [CITE ME], and some models (such as Qwen-7B) may be more sensitive than others.

We also executed a similar ablation experiment as the one above, but instead of testing five different prompts, the models were given the original instruction prompt five times. The pairwise agreement scores for each run can be found in Table 12. It appears that our attempts to restrict stochasticity within the experiment (e.g., the standardization of python, numpy and pytorch seeds), as well as the

¹⁵Schroeder et al. (2024) feature F1 scores in the 0.7-0.9 range, significantly outperforming our own models. However, we were not able to replicate these results given their EncOn model setup. While the authors acknowledged our requests for clarifications in their training and evaluation setup, we have not received any.

Variant A	Variant B	OLMo-32B	OLMo-7B	Qwen-7B	Llama-8B
cautious1	cautious2	0.888	0.972	0.966	0.972
cautious1	concise	0.771	0.808	0.756	0.776
cautious1	descriptive	0.739	0.760	0.587	0.741
cautious1	role-play	0.799	0.689	0.775	0.807
cautious2	concise	0.898	0.807	0.760	0.767
cautious2	descriptive	0.797	0.756	0.585	0.745
cautious2	role-play	0.790	0.695	0.790	0.794
concise	descriptive	0.869	0.859	0.816	0.737
concise	role-play	0.830	0.785	0.598	0.838
descriptive	role-play	0.799	0.815	0.423	0.744

Table 11: Pairwise similarity scores across prompt variants for different language models.

nature of the task, lead to deterministic outputs by the LLMs on the evaluated sample.

F.3 One model to rule them all?

Since comments in the written and oral datasets are fundamentally different (§B.3), we would be tempted to train different models for each. This decision is not obvious however, since it is likely that the model can internally distinguish between oral and written speech. If that occurs, the differences between the texts should not negatively affect the model; in fact, training the model on all datasets could allow it to transfer knowledge between the two text modalities. Table 13 shows the performance of the same model when trained only on oral and written datasets, compared to when trained on the all datasets at once. We find that the models perform better when trained on different splits.

F.4 Decision thresholds

An advantage of using [EncOn](#) classifiers instead of LLMs is that we can quantitatively alter the performance of the classifier by tuning the thresholds required for a comment to trigger an intervention. By making the model more conservative (increasing the threshold), we could hope that we can avoid unnecessary interventions, while making it less conservative (decreasing the threshold) we could make it more proactive.

Fig. 9 shows that the classifier has an inherently low precision score which is largely invariant to the selected threshold. On the other hand, recall falls off rapidly, which incentivises us to use a more liberal approach.

F.5 Is facilitator behavior dependent on domain?

Fig. 10 shows the ratio of facilitative interventions for human and LLM survey participants.¹⁶ Human facilitators are generally invariant to the discussion domain. However, most LLMs tend to under-facilitate in non-forum domains, and over-facilitate in [WT](#), the only dataset featuring a traditional forum structure. A few over-facilitate in general, as seen in Table 4.

Why do we observe this discrepancy between domains? It could be caused by some dataset-specific linguistic markers or latent pattern systematically influencing the judgments of LLM facilitators. Indeed, [WT](#) is the only dataset considered in this study specifically containing escalating discussions, which likely triggers LLM interventions. Human facilitators do not seem to be influenced by the escalating discussions however (Fig. 10). This discrepancy could be attributed to LLMs being hypersensitive to unsafe speech because of their extensive alignment procedures.

An alternative explanation would posit that LLMs may not understand the existence of facilitative interventions in domains other than traditional online forums. Under this assumption the LLMs could abstain from intervening not because they “calculate” that a facilitation intervention would be unnecessary, but rather because they have not been exposed to facilitative interventions during their training. In any case, this assumption does not seem to be hold; Table 14 shows that the LLMs are reasonably competent at detecting facilitators in multiple domains.

¹⁶Technically, LLMs do not take surveys, they execute a task phrased as a survey for humans.

Run A	Run B	OLMo-32B	OLMo-7B	Qwen-7B	Llama-8B
1	2	1.000	1.000	1.000	1.000
1	3	1.000	1.000	1.000	1.000
1	4	1.000	1.000	1.000	1.000
1	5	1.000	1.000	1.000	1.000
2	3	1.000	1.000	1.000	1.000
2	4	1.000	1.000	1.000	1.000
2	5	1.000	1.000	1.000	1.000
3	4	1.000	1.000	1.000	1.000
3	5	1.000	1.000	1.000	1.000
4	5	1.000	1.000	1.000	1.000

Table 12: Pairwise similarity scores across repeated runs for different language models using the original survey instructions.




Figure 9: Classifier performance (**left**: written datasets, **right**: oral datasets) when tuning the decision threshold on the test-set. The higher the threshold, the more conservative the classifier is encouraged to become when predicting a facilitator intervention.




Figure 10: Percentage of times survey participants decided to intervene in discussions by dataset. Some LLMs always over-intervene, while other under-intervene in most domains, except for the traditional forum-structured dataset.

G Survey guidelines

Survey Guidelines: Type of Reinforcement Needed in Comment Threads

Purpose of the Task

These guidelines explain how to rate comments in conversation threads using three parallel rating

columns:

- Positive Reinforcement (1–5)
- Negative Reinforcement (1–5)
- No Reinforcement Needed (1–5)

Your job is to read each comment in its conversational context and assign a strength/certainty rating (1–5) for each category, reflecting how that category applies. All fields need to be filled and the categories are not necessarily mutually exclusive.

Even though your intervention will be delivered at the end of the thread, the reinforcement can refer to any point in the conversation, not just at the end. It can refer to something said earlier, not just the last comment.

You are not being asked whether that conversation correlates to positive or negative reinforcement, but rather whether you, personally, would have used positive or negative reinforcement at a given point. Keep in mind that too many interventions on the moderator’s part may result in an adverse reaction from the participants; no one wants to be

Split	Dataset	Precision	Recall	F1
Written	CeRI	0.320	0.490	0.387
	Wikitactics	0.407	0.717	0.520
oral	Fora	0.409	0.631	0.497
	IQ2	0.399	0.634	0.490
	WHoW	0.418	0.626	0.502
All	CeRI	0.292	0.172	0.216
	Fora	0.410	0.664	0.507
	IQ2	0.406	0.623	0.491
	WHoW	0.422	0.642	0.509
	Wikitactics	0.456	0.364	0.405

Table 13: Summary of classifier performance depending on the training dataset splits on the facilitator prediction task. We compare models trained on all available data with models trained separately on written and oral datasets. Performance is generally similar.

constantly told what to say or how to say it.

Important Notes

- You must add the rationale behind your response in the rational/comments field when:
 - your positive score > 3 , or
 - your negative score > 3 , or
 - you encounter a particular case (e.g., an instance which could fit both negative and positive reinforcement categories).
- If you encounter data that is malformed or noisy to the point that it renders comprehension difficult, flag the instance using the malformed field.

What Each Label Means

Positive Reinforcement

Positive reinforcement is defined as an intervention whose purpose is to encourage a specific behavior. For example, a moderator may ask people to participate or ask a person to elaborate. Positive reinforcement is not related to positive speech; the moderator may be strict or even rude when attempting to encourage specific behaviors (e.g., “I will not tolerate any further discussion on this topic. It’s been decided and we’re moving on.”).

Rate this when a conversation demonstrates behavior that should be encouraged or reinforced in a synchronous exchange. Your task is not simply to identify positive behaviors, but to identify moments where reinforcing them would help sustain or build constructive conversational momentum.

This includes, but is not limited to:

- Helpful or clarifying explanations

- Respectful disagreement
- Constructive feedback
- Polite, empathetic, or supportive tone
- Efforts to de-escalate conflict
- Any contribution that meaningfully improves the quality or direction of the conversation

Negative Reinforcement

The purpose of negative reinforcement is to discourage unwanted or disruptive behaviors. This is correlated, but not necessarily closely related to threats or enforcement of disciplinary action. Negative reinforcement is not related to negative speech; the moderator may be kind and diplomatic when discouraging behaviors (e.g., “I’ve received multiple complaints about your behavior on this forum. Please refrain from making personal attacks and inflammatory comments.”).

Rate this when the comment reflects harmful, disruptive, or inappropriate behavior that should be discouraged, as well as cases where some participants dominate the conversation, or when the conversation starts to go off topic. Other examples include:

- Hostile, dismissive, or disrespectful language
- Personal attacks or insults
- Repeated interruptions or monopolizing the exchange
- Escalating conflict rather than resolving it
- Contributions that reduce the overall quality of the conversation

No Reinforcement Needed

Rate this when the comment is neutral, unremarkable, or does not affect the tone or direction of the conversation. This includes:

- Basic factual statements
- Clarifying questions
- Minimal or vague comments
- Harmless but non-notable contributions

How to Make Decisions

Look at Intent and Impact

Consider what the commenter seems to be trying to do, and how their words might affect others in the thread.

Use the Context

Always check the surrounding messages. A comment that looks fine on its own might escalate tension when viewed in context—and vice versa.

Judge Behavior, Not Opinions

The label should never depend on whether you personally agree with the comment. Focus on tone, clarity, constructiveness, and how it influences the conversation.

1549	Remain Consistent		
1550	Apply the same standards across all comments.	this topic. It's been decided and we're moving	1597
1551	Avoid guessing or over-interpreting; use what is	on.”).	1598
1552	present in the text.	Rate this when a conversation demonstrates be-	1599
		havior that should be encouraged or reinforced in	1600
1553	H LLM Prompts	a synchronous exchange. Your task is not simply	1601
		to identify positive behaviors, but to identify mo-	1602
1554	H.1 Survey Instructions	ments where reinforcing them would help sustain	1603
		or build constructive conversational momentum.	1604
1555	H.1.1 Yes/no Intervention	This includes, but is not limited to:	1605
1556	You are an expert specializing in online moderation	• Helpful or clarifying explanations	1606
1557	and facilitation. You will be presented with a com-	• Respectful disagreement	1607
1558	ment and the comments made before it. Choose	• Constructive feedback	1608
1559	whether you decide to intervene at this point in the	• Polite, empathetic, or supportive tone	1609
1560	discussion.	• Efforts to de-escalate conflict	1610
1561	Key Rules:	• Any contribution that meaningfully improves	1611
1562	Look at Intent and Impact Consider what the	the quality or direction of the conversation	1612
1563	commenter seems to be trying to do, and how their	Negative Reinforcement	1613
1564	words might affect others in the thread.	The purpose of negative reinforcement is to dis-	1614
1565	Use the Context Always check the surround-	courage unwanted or disruptive behaviors. This	1615
1566	ing messages. A comment that looks fine on its	is correlated, but not necessarily closely related to	1616
1567	own might escalate tension when viewed in con-	threats or enforcement of disciplinary action.	1617
1568	text—and vice versa.	Negative reinforcement is not related to negative	1618
1569	Judge Behavior, Not Opinions The label should	speech; the moderator may be kind and diplomatic	1619
1570	never depend on whether you personally agree with	when discouraging behaviors (e.g., “I’ve received	1620
1571	the comment. Focus on tone, clarity, constructiv-	multiple complaints about your behavior on this	1621
1572	ness, and how it influences the conversation.	forum. Please refrain from making personal attacks	1622
1573	Remain Consistent Apply the same standards	and inflammatory comments.”).	1623
1574	across all comments. Avoid guessing or over-	Rate this when the comment reflects harmful,	1624
1575	interpreting; use what is present in the text.	disruptive, or inappropriate behavior that should	1625
1576	Output Respond whether you would intervene	be discouraged, as well as cases where some par-	1626
1577	in this discussion as a facilitator on a scale of 1-5,	ticipants dominate the conversation, or when the	1627
1578	where 1 is “I’m sure I don’t need to intervene”, and	conversation starts to go off topic.	1628
1579	5 is “I’m sure I need to intervene”	Other examples include:	1629
		• Hostile, dismissive, or disrespectful language	1630
1580	H.1.2 Positive vs. Negative reinforcement	• Personal attacks or insults	1631
1581	You are an expert facilitator. You will be presented	• Repeated interruptions or monopolizing the	1632
1582	with a discussion. Decide whether you would write	exchange	1633
1583	a comment intervening after the last comment in	• Escalating conflict rather than resolving it	1634
1584	the discussion.	• Contributions that reduce the overall quality	1635
1585	The Core Rating Task	of the conversation	1636
1586	For every comment, you must write out a re-	No Reinforcement	1637
1587	sponse of 1 (uncertain) to 5 (certain):	Rate this when the comment is neutral, unre-	1638
1588	Positive Reinforcement	markable, or does not affect the tone or direction	1639
1589	Defined as an intervention whose purpose is to	of the conversation.	1640
1590	encourage a specific behavior. For example, a mod-	This includes:	1641
1591	erator may ask people to participate, ask a person	• Basic factual statements	1642
1592	to elaborate etc.	• Clarifying questions	1643
1593	Positive reinforcement is not related to positive	• Minimal or vague comments	1644
1594	speech; the moderator may be strict or even rude	• Harmless but non-notable contributions	1645
1595	when attempting to encourage specific behaviors	Key Rules:	1646
1596	(e.g., “I will not tolerate any further discussion on	Look at Intent and Impact	1647
		Consider what the commenter seems to be trying	1648

1649	to do, and how their words might affect others in		
1650	the thread.		
1651	Use the Context		
1652	Always check the surrounding messages. A com-		
1653	ment that looks fine on its own might escalate ten-		
1654	sion when viewed in context—and vice versa.		
1655	Judge Behavior, Not Opinions		
1656	The label should never depend on whether you		
1657	personally agree with the comment. Focus on tone,		
1658	clarity, constructiveness, and how it influences the		
1659	conversation.		
1660	Remain Consistent		
1661	Apply the same standards across all comments.		
1662	Avoid guessing or over-interpreting; use what is		
1663	present in the text.		
1664	Output		
1665	Respond whether you would intervene in this		
1666	discussion as a facilitator.		
1667	Always respond with the following format and		
1668	nothing else:		
1669	Positive: 1-5		
1670	Negative: 1-5		
1671	No Reinforcement: 1-5		
1672			
1673	H.1.3 Single-Comment Facilitation Prediction		
1674	You are an expert moderator/facilitator. You will		
1675	be presented with a comment (TGT) and the com-		
1676	ments made before it (CTX, one for each comment		
1677	in turn). Choose whether you decide to intervene		
1678	at this point in the discussion.		
1679	Key Rules:		
1680	Look at Intent and Impact Consider what the		
1681	commenter seems to be trying to do, and how their		
1682	words might affect others in the thread.		
1683	Use the Context Always check the surround-		
1684	ing messages. A comment that looks fine on its		
1685	own might escalate tension when viewed in con-		
1686	text—and vice versa.		
1687	Judge Behavior, Not Opinions The label should		
1688	never depend on whether you personally agree with		
1689	the comment. Focus on tone, clarity, constructiveness,		
1690	and how it influences the conversation.		
1691	Remain Consistent Apply the same standards		
1692	across all comments. Avoid guessing or over-		
1693	interpreting; use what is present in the text.		
1694	Output Respond whether you would intervene		
1695	in this discussion as a facilitator. Write 0 if you		
1696	would not intervene, and 1 if you would. Do not		
1697	write anything else, only the number.		
	H.1.4 Single-Comment Facilitation Detection		1698
	You are an expert moderator/facilitator. You will		1699
	be presented with a comment (TGT) and the com-		1700
	ments made before it (CTX, one for each comment		1701
	in turn). Choose whether the comment you are		1702
	reading (TGT) is a comment made by a modera-		1703
	tor/facilitator.		1704
	Key Rules for:		1705
	Look at Intent and Impact Consider what the		1706
	commenter seems to be trying to do, and how their		1707
	words might affect others in the thread.		1708
	Use the Context Always check the surround-		1709
	ing messages. A comment that looks fine on its		1710
	own might escalate tension when viewed in con-		1711
	text—and vice versa.		1712
	Judge Behavior, Not Opinions The label should		1713
	never depend on whether you personally agree with		1714
	the comment. Focus on tone, clarity, constructive-		1715
	ness, and how it influences the conversation.		1716
	Remain Consistent Apply the same standards		1717
	across all comments. Avoid guessing or over-		1718
	interpreting; use what is present in the text.		1719
	Output Respond whether the comment you are		1720
	reading (TGT) is a comment made by a modera-		1721
	tor/facilitator. Write 0 if it was not made by a		1722
	facilitator, and 1 if it was made by a facilitator. Do		1723
	not write anything else, only the number.		1724
	H.2 Ablation prompts		1725
	H.2.1 Cautious1		1726
	You are an expert facilitator. You will be presented		1727
	with a discussion. Decide whether you would write		1728
	a comment intervening after the last comment in		1729
	the discussion.		1730
	The Core Rating Task For every discussion,		1731
	produce three ratings from 1 (uncertain) to 5 (cer-		1732
	tain).		1733
	Positive Reinforcement		1734
	Defined as an intervention whose purpose is to		1735
	encourage a specific behavior. For example, a facil-		1736
	itator may ask people to participate, ask a person		1737
	to elaborate, or encourage constructive discussion.		1738
	Positive reinforcement is not related to positive sen-		1739
	timent; the facilitator may be strict or direct when		1740
	encouraging specific behaviors (e.g., “I will not		1741
	tolerate any further discussion on this topic. It’s		1742
	been decided and we’re moving on.”).		1743
	Assign this rating when the conversation demon-		1744
	strates behavior that should be encouraged or re-		1745
	inforced in a synchronous exchange. Your task is		1746
	not simply to identify positive behaviors, but to		1747

1748	identify moments where reinforcing them would	the thread.	1799
1749	help sustain or build constructive conversational	Use the Context	1800
1750	momentum.	Always consider the surrounding discussion. A	1801
1751	Examples include, but are not limited to:	comment that appears appropriate in isolation may	1802
1752	• Helpful or clarifying explanations	escalate tension in context, and vice versa.	1803
1753	• Respectful disagreement	Judge Behavior, Not Opinions	1804
1754	• Constructive feedback	The labels should never depend on whether you	1805
1755	• Polite, empathetic, or supportive tone	personally agree with the comment. Focus on tone,	1806
1756	• Efforts to de-escalate conflict	clarity, constructiveness, and how the discussion	1807
1757	• Contributions that meaningfully improve the	evolves.	1808
1758	quality or direction of the conversation	Remain Consistent	1809
1759	Negative Reinforcement	Apply the same standards across all discussions.	1810
1760	The purpose of negative reinforcement is to dis-	Avoid guessing or over-interpreting; evaluate only	1811
1761	courage unwanted or disruptive behaviors. This	what is present in the conversation.	1812
1762	is correlated with, but not limited to, disciplinary	Output Respond only using the following for-	1813
1763	action. Negative reinforcement is not related to	mat and nothing else:	1814
1764	negative sentiment; the facilitator may be kind	Positive: 1-5	1815
1765	and diplomatic when discouraging behaviors (e.g.,	Negative: 1-5	1816
1766	“I’ve received multiple complaints about your be-	No Reinforcement: 1-5	1817
1767	havior on this forum. Please refrain from making		1818
1768	personal attacks and inflammatory comments.”).	H.2.2 Cautious2	1819
1769	Assign this rating when the conversation reflects	You are an expert facilitator. You will be presented	1820
1770	harmful, disruptive, or inappropriate behavior that	with a discussion. Decide whether you would write	1821
1771	should be discouraged, as well as cases where par-	a comment intervening after the last comment in	1822
1772	ticipants dominate the conversation or the discus-	the discussion.	1823
1773	sion starts to drift off topic.	The Core Rating Task For every discussion,	1824
1774	Examples include:	provide three ratings on a scale from 1 (uncertain)	1825
1775	• Hostile, dismissive, or disrespectful language	to 5 (certain).	1826
1776	• Personal attacks or insults	Positive Reinforcement	1827
1777	• Repeated interruptions or monopolizing the	Positive reinforcement refers to an intervention	1828
1778	exchange	whose purpose is to encourage a specific behav-	1829
1779	• Escalating conflict rather than resolving it	ior. For example, a facilitator may encourage par-	1830
1780	• Contributions that reduce the overall quality	ticipation, ask someone to elaborate, or reinforce	1831
1781	of the conversation	constructive conversational habits. Positive rein-	1832
1782	No Reinforcement	forcement is not synonymous with positive senti-	1833
1783	Assign this rating when the conversation is neu-	ment; the facilitator may be strict or direct when	1834
1784	tral, unremarkable, or does not meaningfully affect	encouraging desired behaviors (e.g., “I will not tol-	1835
1785	the tone or direction of the discussion. This in-	erate any further discussion on this topic. It’s been	1836
1786	cludes:	decided and we’re moving on.”).	1837
1787	• Basic factual statements	Assign this rating when the discussion demon-	1838
1788	• Clarifying questions	strates behavior that should be encouraged or rein-	1839
1789	• Minimal or vague comments	forced in a synchronous exchange. Do not simply	1840
1790	• Harmless but otherwise unremarkable contri-	identify positive behaviors; instead, identify situa-	1841
1791	butions	tions where reinforcing them would help sustain or	1842
1792	Remember that the role of the facilitator is sup-	strengthen constructive conversational momentum.	1843
1793	portive rather than dominant. While intervention	Examples include, but are not limited to:	1844
1794	is sometimes necessary, unnecessary or frequent	• Helpful or clarifying explanations	1845
1795	intervention degrades the user experience.	• Respectful disagreement	1846
1796	Key Rules Look at Intent and Impact	• Constructive feedback	1847
1797	Consider what the commenter seems to be trying		
1798	to do, and how their words might affect others in		

1848	• Polite, empathetic, or supportive communication	appropriate in isolation may escalate tension in context, and vice versa.	1899
1849			1900
1850	• Attempts to de-escalate conflict	Judge Behavior, Not Opinions	1901
1851	• Contributions that meaningfully improve the quality or direction of the discussion	Your ratings should never depend on whether you personally agree with the content of a comment. Focus instead on tone, clarity, constructiveness, and the comment's effect on the conversation.	1902
1852			1903
1853	Negative Reinforcement		1904
1854	Negative reinforcement refers to an intervention whose purpose is to discourage unwanted or disruptive behaviors. It may involve setting boundaries or discouraging problematic conduct, but does not necessarily involve threats or disciplinary action. Likewise, it is not synonymous with negative sentiment; a facilitator may politely discourage undesirable behavior (e.g., "I've received multiple complaints about your behavior on this forum. Please refrain from making personal attacks and inflammatory comments.").	Remain Consistent	1905
1855		Apply the same standards across all discussions. Avoid over-interpreting intent or inferring information that is not present.	1906
1856			1907
1857			1908
1858			1909
1859		Output Respond <i>only</i> with the following format:	1910
1860			1911
1861		Positive: 1-5	1912
1862		Negative: 1-5	1913
1863		No Reinforcement: 1-5	1914
1864			1915
1865	Assign this rating when the discussion contains harmful, disruptive, or inappropriate behavior that should be discouraged, including situations where participants dominate the conversation or the discussion begins to drift off topic.	H.2.3 Concise	1916
1866	Examples include:	You are an expert facilitator. You will be presented with a discussion. Decide whether you would write a comment intervening after the last comment in the discussion.	1917
1867	• Hostile, dismissive, or disrespectful language		1918
1868	• Personal attacks or insults		1919
1869	• Repeated interruptions or monopolizing the discussion		1920
1870	• Escalating conflict rather than resolving it	The Core Rating Task For every discussion, provide three ratings on a scale from 1 (low intensity) to 5 (high intensity).	1921
1871	• Contributions that reduce the overall quality of the conversation	Positive Reinforcement	1922
1872		Rate highly if the discussion demonstrates constructive behaviors that should be encouraged or reinforced. Examples include:	1923
1873		• Helpful explanations or elaboration	1924
1874		• Respectful disagreement	1925
1875		• Constructive feedback	1926
1876		• Polite, empathetic, or supportive communication	1927
1877		• Attempts to de-escalate conflict	1928
1878	No Reinforcement	• Contributions that improve the quality or direction of the discussion	1929
1879	Assign this rating when the discussion is neutral, unremarkable, or does not meaningfully affect the tone or direction of the conversation. This includes:	This rating reflects opportunities where facilitator reinforcement would strengthen productive conversational dynamics.	1930
1880	• Basic factual statements	Negative Reinforcement	1931
1881	• Clarifying questions	Rate highly if the discussion contains behaviors that should be discouraged. Examples include:	1932
1882	• Minimal or vague comments	• Hostile, dismissive, or disrespectful language	1933
1883	• Harmless but otherwise unremarkable contributions	• Personal attacks or insults	1934
1884		• Repeated interruptions or monopolizing the discussion	1935
1885		• Escalating conflict rather than resolving it	1936
1886		• Going substantially off topic	1937
1887	Your intervention ratings should prioritize surgical precision . Over-moderation is counterproductive, and unnecessary interventions will degrade the participant experience.		1938
1888	Key Rules		1939
1889	Look at Intent and Impact		1940
1890	Consider what each commenter appears to be trying to accomplish, and how their contribution is likely to affect the discussion.		1941
1891	Use the Context		1942
1892	Always evaluate comments in the context of the surrounding discussion. A comment that appears		1943
1893			1944
1894			1945
1895			1946
1896			1947

1948	Contributions that reduce the overall quality of the discussion	1996
1949		1997
1950	This rating reflects opportunities where facilitator intervention would discourage disruptive conversational dynamics.	1998
1951		1999
1952		2000
1953	No Reinforcement	2001
1954	Assign a high rating when the discussion is largely neutral and does not warrant intervention. Examples include:	2002
1955	Basic factual statements	2003
1956		2004
1957	Clarifying questions	2005
1958	Minimal or routine comments	2006
1959	Harmless contributions with no meaningful impact on the discussion	2007
1960		2008
1961		2009
1962	Core Directives Behavior is Key	2010
1963	Evaluate the style, intent, and conversational impact of the comments, never whether you agree with their opinions or conclusions.	2011
1964		2012
1965	Contextual Read	2013
1966	Consider the discussion as a whole. The surrounding context determines whether a contribution is constructive, disruptive, or neutral.	2014
1967		2015
1968	Consistency	2016
1969	Apply the same evaluation criteria uniformly across all discussions. Avoid inferring intentions or information that are not present in the text.	2017
1970		2018
1971		2019
1972		2020
1973		2021
1974	Output Respond using <i>only</i> the following format:	2022
1975		2023
1976	Positive: 1-5	2024
1977	Negative: 1-5	2025
1978	No Reinforcement: 1-5	2026
1979		2027
1980		2028
1981	H.2.4 Descriptive	2029
1982	You are an elite, impartial discussion facilitator tasked with maintaining high conversational quality. You will be presented with a discussion segment. Decide whether an intervention is warranted after the final comment.	2030
1983		2031
1984		2032
1985		2033
1986		2034
1987		2035
1988		2036
1989		2037
1990	The Core Rating Task For each discussion, assign a rating from 1 (low confidence) to 5 (high confidence) for each of the following intervention types.	2038
1991		2039
1992	Positive Reinforcement (Encouraging Growth)	2040
1993	Assign a higher rating when the discussion demonstrates behaviors that actively improve constructive dialogue and would benefit from reinforcement. Examples include:	2041
1994		2042
1995		2043
		2044
		2045
		2046

2047	Evaluate conversational behavior only. Ratings		
2048	must never depend on agreement or disagreement		
2049	with the viewpoints expressed.		
2050	Consistency		
2051	Apply these criteria uniformly across all discus-		
2052	sions. Avoid inferring motives or information that		
2053	are not supported by the text.		
2054	Output Respond <i>only</i> using the following for-		
2055	mat:		
2056	Positive: 1-5		
2057	Negative: 1-5		
2058	No Reinforcement: 1-5		
2059			
2060	H.2.5 Role-play		
2061	You are assuming the role of a highly skilled, proac-		
2062	tive discussion facilitator. Your primary objective		
2063	is to determine whether an intervention should oc-		
2064	cure after the final comment based on the preceding		
2065	discussion.		
2066	The Core Assessment Matrix For the discus-		
2067	sion, assign a score from 1 (low confidence) to 5		
2068	(high confidence) for each of the following inter-		
2069	vention types.		
2070	Positive Encouragement (Boosting Momen-		
2071	tum)		
2072	Assign a higher score when the discussion con-		
2073	tains constructive behaviors that would benefit		
2074	from reinforcement. Consider whether encour-		
2075	aging these behaviors would improve the quality,		
2076	depth, or collaboration of the discussion.		
2077	Examples include:		
2078	• Helpful or insightful explanations		
2079	• Respectful and well-reasoned disagreement		
2080	• Constructive or actionable feedback		
2081	• Empathetic or supportive engagement		
2082	• Attempts to de-escalate conflict		
2083	• Contributions that meaningfully improve the		
2084	discussion		
2085	Positive encouragement reinforces desirable con-		
2086	versational behaviors rather than agreement with		
2087	the opinions expressed.		
2088	Corrective Guidance (Steering Away from Pit-		
2089	falls)		
2090	Assign a higher score when the discussion ex-		
2091	hibits behaviors that should be discouraged to pre-		
2092	serve conversational quality.		
2093	Examples include:		
2094	• Hostile, dismissive, or disrespectful commu-		
2095	nication		
	• Personal attacks or inflammatory remarks		2096
	• Off-topic tangents or repetitive, low-value		2097
	contributions		2098
	• Dominating the discussion or discouraging		2099
	participation		2100
	• Escalating conflict without moving toward res-		2101
	olution		2102
	Corrective guidance targets disruptive conversa-		2103
	tional behaviors rather than disciplinary action or		2104
	disagreement with viewpoints.		2105
	Inaction Required (Status Quo)		2106
	Assign a higher score when no facilitator inter-		2107
	vention appears necessary. Typical examples in-		2108
	clude:		2109
	• Neutral factual statements		2110
	• Requests for clarification		2111
	• Routine conversational exchanges		2112
	• Contributions that have little effect on the dis-		2113
	cussion's direction or tone		2114
	Rules of Engagement Observe Dynamics		2115
	Evaluate the discussion as a whole before assign-		2116
	ing ratings. The surrounding context determines		2117
	whether intervention would improve the conversa-		2118
	tion.		2119
	Behavior Over Belief		2120
	Assess only the conversational behavior, includ-		2121
	ing tone, constructiveness, and impact. Never eval-		2122
	uate the underlying opinions or viewpoints.		2123
	Pragmatism		2124
	The ratings should reflect whether facilitator in-		2125
	tervention would meaningfully improve the discus-		2126
	sion. Avoid unnecessary intervention when the		2127
	discussion is progressing appropriately.		2128
	Output Respond <i>only</i> using the following for-		2129
	mat:		2130
	Positive: 1-5		2131
	Negative: 1-5		2132
	No Reinforcement: 1-5		2133
			2134

Model	Pr.	Rec.	$F1_p$	$F1_n$	Sup.
CeRI					
LLaMa70B	0.641	0.949	0.765	0.903	335
LLaMa8B	0.550	0.899	0.683	0.857	335
ModBert	0.571	0.763	0.654	0.890	152
OLMo32B	0.845	0.899	0.871	0.959	335
OLMo7B	0.345	0.367	0.356	0.793	335
Qwen32B	0.910	0.899	0.904	0.971	335
Qwen7B	0.828	0.608	0.701	0.923	335
All	0.670	0.769	0.705	0.899	335
WT					
LLaMa70B	0.410	0.641	0.500	0.500	336
LLaMa8B	0.413	0.489	0.448	0.591	336
ModBert	0.509	0.780	0.616	0.675	305
OLMo32B	0.413	0.382	0.397	0.638	336
OLMo7B	0.342	0.588	0.433	0.361	336
Qwen32B	0.509	0.214	0.301	0.733	336
Qwen7B	0.476	0.153	0.231	0.733	336
All	0.439	0.464	0.418	0.604	336
Fora					
LLaMa70B	0.571	0.830	0.677	0.753	382
LLaMa8B	0.568	0.741	0.643	0.755	382
ModBert	0.469	0.618	0.533	0.687	2736
OLMo32B	0.738	0.563	0.639	0.837	382
OLMo7B	0.215	0.104	0.140	0.695	382
Qwen32B	0.824	0.556	0.664	0.859	382
Qwen7B	0.703	0.333	0.452	0.807	382
All	0.584	0.535	0.535	0.770	382
IQ2					
LLaMa70B	0.418	0.774	0.543	0.643	374
LLaMa8B	0.394	0.617	0.481	0.662	374
ModBert	0.550	0.681	0.608	0.724	2515
OLMo32B	0.709	0.487	0.577	0.852	374
OLMo7B	0.256	0.174	0.207	0.724	374
Qwen32B	0.824	0.530	0.646	0.880	374
Qwen7B	0.826	0.165	0.275	0.836	374
All	0.568	0.490	0.477	0.760	374
WHoW					
LLaMa70B	0.476	0.901	0.623	0.651	364
LLaMa8B	0.449	0.661	0.535	0.676	364
ModBert	0.507	0.617	0.557	0.696	1869
OLMo32B	0.699	0.479	0.569	0.832	364
OLMo7B	0.274	0.140	0.186	0.727	364
Qwen32B	0.814	0.471	0.597	0.857	364
Qwen7B	0.733	0.182	0.291	0.815	364
All	0.565	0.493	0.480	0.751	364
All datasets					
LLaMa70B	0.419	0.682	0.518	0.740	2000
LLaMa8B	0.396	0.568	0.465	0.756	2000
ModBert	0.521	0.692	0.594	0.734	7577
OLMo32B	0.567	0.468	0.509	0.852	2000
OLMo7B	0.239	0.229	0.220	0.714	2000
Qwen32B	0.647	0.445	0.519	0.883	2000
Qwen7B	0.594	0.240	0.325	0.852	2000
All	0.483	0.475	0.450	0.790	2000

Table 14: Classifier performance for facilitator detection using positive-class precision, recall and F1 ($F1_p$), as well as negative-class F1 ($F1_n$). Rows are struck through when either $F1_p$ or $F1_n$ collapses (< 0.15). Best non-collapsed $F1_p$ and $F1_n$ values for each dataset split are highlighted in bold. “All” rows report macro averages across models.