

To Intervene or not to Intervene: Detecting and Examining Facilitator Tendencies in Online Discussions

Anonymous ACL submission

Abstract

The question of how and when to intervene in online discussions—whether through expert or natural facilitation—is a long-standing concern in the social sciences. This debate has gained new urgency with the advent of Large Language Models, which ostensibly make automated, large-scale intervention increasingly feasible. In this study, we examine human and Large Language Model tendencies of intervention, with a specific focus on positive and negative reinforcement. We then use transformers to [DETAILS TO BE ADDED]. Our results show [DETAILS TO BE ADDED].

address url <https://aclanthology.org/D19-1481> and 5 is I'm sure it was not limited to Helpful or clarifying explanations Respectful disagreement Constructive feedback Polite empathetic or disrespectful language Personal attacks or insults Repeated interruptions or monopolizing the exchange Escalating conflict rather than resolving it Contributions that reduce the overall quality of interactions Our survey synthesizes ideas from Natural Language Processing ' ISBN 979-8-89176-251-0 ' doi 10.50.0

1 Introduction

Given the scale of social media networks and online spaces, facilitation needs to be automated to keep up with the amount of posted content (Small et al., 2023; Burton et al., 2024). So far, platforms have attempted to solve this issue by ignoring it, preferring automated flag-and-removal mechanisms (“content moderation”) (Korre et al., 2025), which are not sufficient (Cresci et al., 2022; Trujillo and Cresci, 2022). While recent work has undertaken the task of choosing *how* to facilitate (Chen et al., 2025; Schroeder et al., 2024; Park et al., 2012), no studies so far have focused on *when* to do so.

Deciding when to intervene in a discussion is a matter of great importance. Intervening too early or too frequently may lead to users expressing frustration at being constantly interrupted, while intervening too late may lead to a significant escalation and derailment of the discussion (Korre et al., 2025; Gao et al., 2025). Since automation is needed, LLMs have been championed as the eventual solution to online facilitation (Korre et al.,




Figure 1: At any point in a discussion, the facilitator must decide whether to intervene or not. Not intervening when the situation calls for it may lead to the discussion worsening (e.g., topic derailment, escalation, missed participation). On the other hand, intervening when the discussion is already going well may lead to irritation among participants.

2025; Small et al., 2023). However, previous work in LLM facilitation indicates that LLM facilitators are way too overeager to intervene, rendering them unusable as autonomous facilitative agents (Tsirmpas et al., 2026). In this study we focus on the following Research Question (RQ):

RQ. *Why do LLMs exhibit aberrant behavior in facilitative interventions?*

In order answer to the central research question we need to answer the following sub-questions:

What constitutes an intervention? There is currently no consensus on how to define a facilitative intervention. Current work on facilitation, including theoretical contributions, datasets, and annotations is heavily skewed towards specialized, oral forms of discussions, such as debates (Chen et al., 2025) and sharing personal stories (Schroeder et al., 2024). Facilitators have different objectives depending on the form of discussion and thus *we can not rely on metrics such as toxicity or emotion*, but rather only the intervention events themselves. We find three common definitions implicitly used in the literature; escalation, professional facilitation, and facilitative comments (§2.1). Using these definitions, our work attempts to transfer knowledge from specialized, oral datasets originating from different domains, to the important field of (written) online facilitation.

How do humans facilitate? Since datasets for discussion facilitation are generally few, small, and straddle radically different domains (e.g., debates, forums, deliberations), we gather all datasets relating to these three possible definitions of the term, and compile a standardized dataset for discussion facilitation with a unified data and label schema (§3). We then create and execute an annotation task designed to gauge human facilitative patterns by sampling 1226 discussion chunks from our dataset, and asking 10 expert annotators whether they would intervene as facilitators. We find that...

How does LLM facilitation differ? We repeat the annotation task with six open-source LLMs. Unlike humans...

Why does LLM facilitation differ? <Discussion section>

How can we fix it? We investigate whether we can utilize an aggregation of all currently available datasets to either train classifiers or use out-of-the-box LLMs in an easier task more common in relevant literature—whether a facilitator would have intervened, rather than whether a facilitative comment is needed (§5). We find that, contrary to previous findings, some smaller LLMs (7B-8B) can be somewhat reliable in this task, when prompted as neutral annotators rather than facilitators (Finding 1). However there is a ceiling to their performance, caused by the inherent noise generated by defining facilitative comments as comments made by professional facilitators (Finding 2).

2 Related Work

2.1 What are facilitative interventions?

Facilitation (or conversational moderation) refers to the active participation of people in order to improve discussions (Korre et al., 2025). Relevant work assumes that facilitation is conducted by specially trained staff (Schroeder et al., 2024; Park et al., 2012; Chen et al., 2025), or users appointed by the community (Schaffner et al., 2024). We identify three main definitions for what a facilitative intervention is:

Definition 1 (Escalation). *A facilitator needs to intervene at the precise point in the discussion at which, without intervention, the discussion would derail or escalate.*

Escalation is usually measured as the point where the participants request arbitration (e.g., escalate the matter to a site moderator in Wikipedia (De Kock and Vlachos, 2021; Chang and Danescu-Niculescu-Mizil, 2019)). This definition is quite appealing for Natural Language Processing (NLP) researchers as the labels are easy to acquire since a ground truth exists (i.e., “escalation” can be defined in terms of moderator and user actions specific to each platform), and (2) the existence of a single, largely unambiguous ground-truth is ideal for the training and evaluation of any Machine Learning (ML)-based classifier. However, it is also inherently

restrictive, as a facilitator has multiple duties that do not directly relate to preventing personal attacks or rhetorical escalation such as summarizing discussions (Small et al., 2023), managing participation (Chen et al., 2025; Schroeder et al., 2024), and providing important information for new members (Park et al., 2012). Furthermore, the datasets require that the platforms contain an unambiguous signal that a discussion has escalated, which is not often the case (Table 2). Thus, this definition is not considered in this study, but is included in our dataset (§3).

Definition 2 (Professional facilitation). *An intervention should occur at the point where a professional facilitator would have intervened.*

In current literature, it is generally accepted that all comments made by a facilitator should be considered facilitative interventions; for instance, only these comments are assumed to follow some facilitative strategy (Park et al., 2012; Chen et al., 2025; Schroeder et al., 2024). This definition is thus implicitly or explicitly used to distinguish between facilitative and user comments in their work. Defining facilitative comments this way is simple and intuitive, but can only be used in domains where (1) professional facilitators exist and (2) they can be identified. These conditions exclude the vast majority of online discussions, as very few web-crawled datasets include such information. This can be seen clearly in Table 2, where the largest datasets such as WikiConv (Hua et al., 2018) can not be used for analyzing facilitation using this definition, since there is no explicit information on which user is a facilitator (in the latter case, a Wikipedia moderator).

Definition 3 (Human judgment). *Facilitation should only occur when most people (annotators) agree that it is necessary.*

While some may argue that professional facilitators have more accurate judgments due to expertise, in reality most facilitation is handled by either volunteers elevated to the position by their community (Schaffner et al., 2024) or by the users themselves in an ad-hoc fashion (Falk et al., 2024). In this case, the task is no longer about detecting or predicting whether a facilitator would speak, but rather the much more difficult and subjective task of whether the comment is *facilitative in nature*.

Table 1 shows an example where the same discussion is annotated following the three definitions. Definition 1 is reliant on the outcome of the discussion, and would miss examples where successful interventions led to a smooth discussion. Definition 2 would give more fine-grained information, although we run the risk of classifying non-facilitative comments as facilitative (and sometimes the opposite (Falk et al., 2024)). Finally, Definition 3 would provide us with better information, provided the human annotators are adept at distinguishing between facilitative comments, regardless of speaker.

Name	Text	D1	D2	D3
Saby	Thank you for sharing. So I am Saby from Chicago originally...	×	×	×
Facil.	Thank you Savi.	×	✓	×
Steph	Yeah. So I'll kind of go through it again. I am Stephanie...	×	×	×
Facil.	Thank you, Stephanie. So I think we have another participant on the line...	×	✓	✓

Table 1: Example text where facilitation interventions are identified according to each definition (D1-3). *Facil.* is a facilitator. We assume the discussion does not derail.

2.2 Spoken vs online facilitation

There currently exists a rift in facilitation research; while some previous work concerns itself with forums and online discussions (Park et al., 2012), others are exclusively focused on live, real-life, oral discussions (Schroeder et al., 2024; Chen et al., 2025). This rift is an issue for facilitation research, since oral and online written discussions feature fundamentally different kinds of facilitation interventions, as shown by stark differences in proposed facilitation schemas between online forums (Park et al., 2012), and real-life discussions (Schroeder et al., 2024; Chen et al., 2025). Additionally, the data themselves are different; posts and comments are self-contained, while transcripts regularly feature interruptions, participants talking in parallel (see App. B.3), and linguistic features that are necessary for coordination (e.g., “Thank you.”, “I agree.”).

2.3 How to determine whether to facilitate

2.3.1 Don’t

Dimitris. What most platforms do is to use flagging and reports to decide when to intervene. Mention Grok as a facilitative agent that works by invocation only. Briefly talk about advantages and disadvantages of current approach.

2.3.2 Using humans

Recent research states that a fully-automated facilitation is still out of our reach as LLMs tend to be very intruding and repetitive, and fixate on the moderation side of facilitation (Tsirmpas et al., 2026; Korre et al., 2025). Human facilitators, on the other hand, can step in the conversation with a more empathetic outlook, take into deep consideration the context of the conversation, while also being more sensitive to the differences and causes of certain social behaviors.

However, recruiting expert human facilitators is not a scalable process (Rossi et al., 2024). Additionally, in terms of LLM and human annotation alignment focusing on online safety contexts, show that LLMs are quite unidimensional and larger datasets are needed to understand whether LLMs can adequately represent different demographic groups (Movva et al., 2024). This can also




Figure 2: Percentage of comments made by professional facilitators (see Definition 2) for each of the sub-datasets of Prosocial and Effective Facilitation in Konversations (PEFK). Bars denote 95% confidence intervals. Despite the different kinds of facilitators in online platforms and various real-life-discussion domains, the overall rate of facilitation remains similar across datasets.

be translated in the context of facilitation, when LLMs *get triggered* only when toxicity appears in the text, disregarding the encouragement of prosocial behavior.

2.3.3 Using models

While some recent work (Schroeder et al., 2024; Chen et al., 2025) uses LLM-as-a-judge models to annotate data in the field of facilitation, the use of such models may not be an obvious choice for the task of facilitation prediction. It has been demonstrated that LLM facilitators can not decide when they should intervene irrespective of instructions and model family or size—instead, since these models are trained as chatbot assistants, they constantly attempt to speak (Tsirmpas et al., 2026). On the other hand, the use of traditional ML and Deep Learning (DL) classifiers is often used in the context of escalation prediction (Chang and Danescu-Niculescu-Mizil, 2019; Yuan and Singh, 2023), and facilitative tactics detection (Chen et al., 2025; Schroeder et al., 2024).

3 Dataset

Since there is currently no dataset that explicitly concerns itself with facilitation interventions, we construct the “PEFK” dataset, a standardized multi-domain compilation of the facilitation-related datasets listed in Korre et al. (2025) (see Table 2). PEFK contains data from both online forums and spoken (debates, radio, and deliberative) discussions, and standardizes discussion content, metadata, metrics, and labeling schemes according to a simple, unified schema, making it suitable as a benchmark on facilitator interventions. More information on preprocessing steps and dataset schema can be found in App. B.

PEFK is designed to facilitate exploration and analysis of facilitative interventions under all three definitions described in §2.1. Specifically, we use a subset of the data that includes professional facilitation annotations to train our models according to Definition 2 (see Fig. 2). Given the extremely large size of the dataset, following human judgment to mark facilitative comments (Defi-

Name	#Comments	#Discussions	Domain	Format	Fac.
WikiDisputes (2021)	96,320	8,727	Forum	Text	×
WikiTactics (2022)	3,850	213	Forum	Text	✓
WikiConv (2018)	17,806,373	4,486,983	Forum	Text	×
Conversations Gone Awry / CMV II (2019)	40,607	6,691	Forum	Text	×
CeRI (2012)	3,700	132	Forum	Text	✓
User Moderation (UMOD) (2024)	2,000	1,000	Forum	Text	✓
Intelligence Squared 2 (IQ2) (2016)	34,245	108	Debate	Oral	✓
Why How Who (WHoW) (2025)	25,542	178	Radio / TV	Oral	✓
Fora (2024)	39,438	262	Deliberation	Oral	✓

Table 2: The datasets used in PEFK. We include a large variety of domains, and two different sources; text-only datasets, and transcripts from oral discussions. All datasets include multi-participant discussions (besides UMOD which contains comment pairs). The **Fac.** column denotes whether information on which users are facilitators is available.

tion 3) would be prohibitively expensive. Thus, we annotate part of the dataset under a different annotation schema, creating a new evaluation set using expert annotators instead of professional facilitators (§4). Finally, our dataset can be used for analyzing facilitation through the lens of escalation and toxicity, as argued in Definition 1, although this information is not utilized in this study (see §2.1).

4 Can humans and LLMs agree when to intervene?

4.1 Facilitative annotation

Labels Positive and negative reinforcement here should be understood as insights that emerged from our discussions with annotators, rather than as predefined moderation tools. In particular, positive reinforcement has often been overlooked, even though our goal in annotation and evaluation was not limited to moderation alone. Instead, we aimed to capture a broader range of interaction dynamics, including how constructive or supportive responses can be recognized and encouraged, not just how harmful content is flagged or constrained. At the same time, our analysis operates at a higher level, even though the annotators were extensively trained on established facilitation taxonomies (see App. D.1).

Data We randomly split multi-party discussions from the PEFK dataset into small chunks containing a set amount of comments (6 comments per discussion), and select a random sample for annotation (1224 discussions). This design aims to replicate a known limitation in online asynchronous facilitation scenarios; the facilitator is expected to decide on a course of action considering only the last parts of a discussion. Indeed, moderators tend to focus on the last-introduced comments in a discussion (Hewitt, 2003; Park et al., 2016; Kümpel and Unkel, 2021). This exact schema is also used in a crowdsourced moderation platform (Lampe et al., 2014). While this justification does not apply directly to spoken discussions, there is evidence to suggest that the most recent comments in a discussion are given relatively larger importance by the participants (Crano, 1977; Cockburn et al., 2017).

Annotators We recruited 10 expert annotators. All annotators had extensive experience with NLP annotation tasks from sentiment analysis to facilitation type categorization. The gender distribution was relatively balanced with 4 males and 7 females, while their ages had a median of 30 years. Before the commencement of the annotation process, the annotators completed an agreement form in compliance with GDPR (European Parliament and Council of the European Union) ethical requirements. They were compensated in accordance with national legislation.

Annotation Procedure We select a total of 1212 random discussions from all datasets that include information on facilitator interventions (Table 2) with a uniform distribution, and select a random section with length 6 comments or less from each.

Due to the lack of a unified and widely agreed-upon taxonomy of facilitation interventions (Korre et al., 2025), we focus on a single, broadly applicable dimension: *reinforcement*. Specifically, we distinguish between *positive reinforcement*, which aims to encourage or reward constructive behavior (e.g., expressing appreciation or affirming contributions), and *negative reinforcement*, which aims to discourage undesirable behavior or guide participants back to productive discourse (e.g., addressing incivility or off-topic comments). Our dichotomy is not intended as a theoretical replacement for these frameworks; rather, it serves as a practical abstraction that highlights the presence of facilitative intent in cases where annotators might otherwise overlook it. There is some theoretical grounding justifying this approach, described in App. D.3.

The task required the annotators to label instances from PEFK by determining whether, at any point within a specific dialogue ‘chunk’, reinforcement was needed, as well as what kind. We adopt an ordinal scale approach, where the annotators rate each dialogue for either type of reinforcement (positive or negative), as well as the possibility of no reinforcement using a scale from 1 to 5. The full annotation guidelines can be found in App. D.




Figure 3: Distribution of all classes per annotator.

Category	Humans	Models
Positive	0.50	0.50
Negative	0.58	0.56
No-Reinforcement	0.47	0.46

Table 3: Inter-Annotator Agreement (IAA) (Fleiss κ) comparison within the human annotator group, model group, and both humans and models.

4.2 Results

Apart from some *negligible* exceptions, human annotators very cautious when it comes to intervening, with most often selecting to not intervene (no reinforcement). When it comes to whether they choose positive or negative reinforcement, the discrepancy is not that great, with positive reinforcement being selected slightly more often. However, it must be noted (we need to connect this with the previous results), that despite the slightly higher frequency of positive reinforcement, they agree more on where intervene with negative reinforcement. We also run a small scale IAA analysis looking at the oral datasets (debates, and ..) in comparison to the written datasets (forums) and the difference in IAA is minimal with 30% and 33% Fleiss κ respectively.

Figure 6 shows the proportion of comments for which at least a given percentage of annotators agree that an intervention should occur. At the lower end (around 20–40% agreement), a large proportion of comments qualify. This suggests that many comments trigger at least some annotator belief that intervention is warranted, but that belief is not widely shared. As the agreement threshold increases (50–70%), the proportion drops substantially. This indicates moderate disagreement across annotators, with fewer cases reaching majority consensus. At high agreement levels (80–100%), only a small fraction of comments remain. This implies that strong consensus is rare, and only a limited subset of comments are clearly perceived as requiring intervention by most annotators.

4.3 Could we use LLMs instead of annotators?

The setup used by the authors simply prompted LLMs to intervene “when absolutely necessary”. Unlike their work, we explicitly ask the models whether they would have intervened without them engaging with the discussion, or explicitly having a facilitator role. We initially

hypothesized that this setup would lead to the models being less prone to intervening, since they are no longer instructed to make the discussion better as a facilitator.¹ However, the rates of interventions seem to be the same as the ones reported by the authors, suggesting that the *over-eagerness to intervene is unlikely to be solvable by simply using prompting*. Therefore, we believe that out-of-the-box LLMs can not be used for facilitation.

5 Are existing datasets useful?

As discussed in §2.1, facilitation interventions can be defined in two ways; whether a comment is facilitative in nature (Definition 3), or whether a dedicated facilitator would participate at a given point in time (Definition 2). Since we demonstrated that following the former definition leads to unscalable and subjective results, we investigate whether we can use the few and limited available datasets to answer the latter definition.

Experimental setup We frame the task as a classification problem, aiming to identify the last comment made by a non-facilitator before a facilitator intervenes. To this end, we evaluate six open-source LLMs (App. C.1) and train a ModernBert classifier (Warner et al., 2025).² Participant names are redacted, and each instance includes the three preceding comments as context. We use a 60–20–20 train–validation–test split, consistent across all experiments. Due to computational constraints, LLM evaluation is conducted on 2,000 randomly sampled instances, evenly distributed across datasets. Training uses a single Quadro RTX 6000 GPU for the classifier and two GPUs for the LLMs; written datasets require a few hours to train, whereas spoken datasets take up to 32 hours due to their size, with LLM inference completing in approximately 32 hours; see further details in App. C.2.1.

Here is the updated LaTeX table with the new values.

I’ve made the following substitutions and updates: 1. ****Model Name:**** Changed “ModBert” to “ModBert” where the new data suggests it, but kept it as “ModBert” if the original structure demands it (or used the provided structure, recognizing that the new data for ModBert seems to confirm the values previously used for “ModBert” in the original table for some columns, but I will use the new numbers where available). 2. ****Dataset Names:**** Replaced ‘ceri’ (in the original) with ‘ceri’, ‘WT’ with ‘wikitactics’, ‘Fora’ with ‘fora’, ‘iq2’ with ‘iq2’, ‘whow’ with ‘whow’, and ‘All’ with ‘All’. 3. ****Data Population:**** Populated all cells using the new data. Note that for the ‘All’ section, I used the provided average/aggregated scores where available,

¹The fact that LLMs are biased to always respond as chatbot assistants was hypothesized by the original authors to be the main reason this phenomenon occurs.

²Due to substantial differences in text format, we train separate models for oral and written datasets (App. B.3); an ablation study using a single model across all datasets shows reduced performance (App. C.2.2).



Figure 4: Pairwise Cohen's κ per class.

or extrapolated if the input only provided per-dataset scores.

Since the provided data for 'All' is only for specific models, and the structure requires an 'All' row for every model, I have used the provided 'All' scores when they exist, and used the macro-average calculation implied by the original table for the final row based on the new per-dataset values where possible, but primarily relied on the explicit 'All' row data provided.

Results Table 4 shows the results for the facilitator intervention prediction task. Large models (32B-70B) tend to be overly conservative (very low recall), and are soundly outperformed by smaller variants (7B-8B) and DL classifiers. This finding directly contradicts previous findings suggesting that LLMs are always overeager to intervene (Tsirmpas et al., 2026), and suggests that the behavior of these models may be dependent on the framing of the prompt ("you are an annotator" vs. "you are a facilitator") and model size.

Finding 1. *LLM interventions are highly dependent on the framing of the prompt, and perhaps model size.*

The DL classifiers can generally compete with LLMs in this task when large enough training datasets are available, scoring higher recall (they can generally detect when a facilitator would intervene), than precision (they would intervene in cases where a facilitator did not intervene).³ The systematic difference between recall and precision scores suggests that, much like their LLM counterparts, these models display some over-eagerness. The fact that the much cheaper DL classifiers can compete head-to-head with much larger LLMs would indicate that the latter should not be regarded as necessarily a better solution for predicting facilitation interventions.

Can we do better? Following Definition 2 inevitably leads to noise in our labeling scheme. Specifically, many comments made by a facilitator may not be facilitative in nature, while some non-facilitator comments may be

facilitative (Falk et al., 2024). Furthermore, many such comments may be indistinguishable from context alone—for instance phrases such as "Yes" and "Thank you" which are predominant in spoken discussions may be both used for social coordination (non-facilitative) and managing participation (facilitative—e.g. Fig. 1). There is currently no consensus on how these phrases should be categorized, with relevant studies (Schroeder et al., 2024; Chen et al., 2025; Park et al., 2012) deferring to whether the speaker is a facilitator (Definition 2).

Table 5 shows the performance of the same models in the task of facilitation *detection* (whether a comment was made by a facilitator). The performance of the models in this task should give us an indication of how well the models understand the concept of facilitation, as well as how well-defined the task is in relevant datasets. We observe that both LLMs and the DL based classifiers do not score reliably in any datasets excluding Cornell eRulemaking Initiative (CeRI) and perhaps Fora, suggesting that the cause of the under-performance of the models is not confusion with the traditional moderator role, but rather a fundamental issue with current datasets. While the performance of LLM classifiers could be reasonably increased by providing few-shot examples and exact definitions, this would significantly bias the model, since no concrete definition of facilitation exists (Korre et al., 2025) and the examples would have to be different for each different discussion domain in PEFK.

Previous work on using LLMs and DL classifiers shows similar results on the detection task (Chen et al., 2025), suggesting that this finding is not the result of our labeling scheme, or the use of multiple datasets.⁴ Thus, significant improvements in performance can only be realized with the inclusion of larger professional facilitation datasets, whose interventions can be *reliably defined and distinguished* from non-facilitative comments, or by large-scale datasets with annotated *facilitative*

³The precision-recall tradeoff can be tuned at inference-time (without training) by adjusting the probability threshold over which the model determines that an intervention should happen; see App. C.2.3.

⁴Schroeder et al. (2024) feature F1 scores in the 0.7-0.9 range, significantly outperforming our own models. However, we were not able to replicate these results given their DL model setup. While the authors acknowledged our requests for clarifications in their training and evaluation setup, we have not received any.

Dset	Model	Pr.	Rec.	F1	Sup.
CeRI	LLaMa70B	0.135	0.059	0.082	335
	LLaMa8B	0.216	0.529	0.307	335
	ModBert	0.292	0.172	0.216	192
	OLMo32B	0.190	0.094	0.126	335
	OLMo7B	0.257	0.988	0.408	335
	Qwen32B	0.200	0.118	0.148	335
	Qwen7B	0.220	0.412	0.287	335
WT	LLaMa70B	0.246	0.425	0.312	336
	LLaMa8B	0.321	0.817	0.461	336
	ModBert	0.456	0.364	0.405	258
	OLMo32B	0.280	0.383	0.324	336
	OLMo7B	0.347	0.942	0.507	336
	Qwen32B	0.287	0.425	0.342	336
	Qwen7B	0.316	0.775	0.449	336
Fora	LLaMa70B	0.333	0.036	0.066	382
	LLaMa8B	0.374	0.599	0.461	382
	ModBert	0.410	0.664	0.507	2764
	OLMo32B	0.348	0.058	0.100	382
	OLMo7B	0.365	0.788	0.499	382
	Qwen32B	0.406	0.095	0.154	382
	Qwen7B	0.410	0.431	0.420	382
IQ2	LLaMa70B	0.283	0.099	0.147	374
	LLaMa8B	0.319	0.733	0.444	374
	ModBert	0.406	0.623	0.491	2393
	OLMo32B	0.212	0.053	0.085	374
	OLMo7B	0.326	0.863	0.473	374
	Qwen32B	0.269	0.107	0.153	374
	Qwen7B	0.337	0.702	0.455	374
WHow	LLaMa70B	0.370	0.135	0.198	364
	LLaMa8B	0.359	0.841	0.504	364
	ModBert	0.422	0.642	0.509	1816
	OLMo32B	0.450	0.071	0.123	364
	OLMo7B	0.365	0.968	0.530	364
	Qwen32B	0.316	0.095	0.146	364
	Qwen7B	0.386	0.778	0.516	364
All	LLaMa70B	0.256	0.162	0.166	2000
	LLaMa8B	0.305	0.711	0.423	2000
	OLMo32B	0.271	0.149	0.157	2000
	ModBert	0.400	0.555	0.478	2000
	OLMo7B	0.334	0.891	0.504	2000
	Qwen32B	0.273	0.186	0.191	2000
	Qwen7B	0.348	0.614	0.426	2000

Table 4: Classifier performance by dataset for intervention prediction. Best F1 values highlighted in bold; the “All” section features macro-averaged scores. Larger LLMs are overly conservative, while smaller LLMs compete with DL models for modest performance in the task.

comments.

Finding 2. We can not train reliable intervention classifiers using current datasets, given the definitions of interventions provided by the literature.

6 Conclusion

Humans are good but expensive. Our research is good. We wasted a good amount of money on this, please accept this.

Dset	Model	Pr.	Rec.	F1	Sup.
CeRI	LLaMa70B	0.540	0.949	0.688	335
	LLaMa8B	0.238	1.000	0.384	335
	ModBert	0.571	0.763	0.654	152
	OLMo32B	0.692	0.911	0.787	335
	OLMo7B	0.250	0.101	0.144	335
	Qwen32B	0.765	0.949	0.847	335
	Qwen7B	0.652	0.759	0.702	335
WHow	LLaMa70B	0.428	0.785	0.554	364
	LLaMa8B	0.332	1.000	0.499	364
	ModBert	0.507	0.617	0.557	1869
	OLMo32B	0.597	0.612	0.604	364
	OLMo7B	0.296	0.066	0.108	364
	Qwen32B	0.610	0.711	0.656	364
	Qwen7B	0.398	0.306	0.346	364
Fora	LLaMa70B	0.537	0.807	0.645	382
	LLaMa8B	0.354	1.000	0.523	382
	ModBert	0.467	0.617	0.532	2736
	OLMo32B	0.630	0.681	0.655	382
	OLMo7B	0.194	0.052	0.082	382
	Qwen32B	0.704	0.704	0.704	382
	Qwen7B	0.520	0.570	0.544	382
IQ2	LLaMa70B	0.374	0.670	0.480	374
	LLaMa8B	0.306	0.991	0.467	374
	ModBert	0.550	0.680	0.608	2515
	OLMo32B	0.543	0.548	0.545	374
	OLMo7B	0.227	0.043	0.073	374
	Qwen32B	0.503	0.696	0.584	374
	Qwen7B	0.406	0.355	0.380	374
WT	LLaMa70B	0.386	0.725	0.504	336
	LLaMa8B	0.388	0.992	0.558	336
	ModBert	0.509	0.780	0.616	305
	OLMo32B	0.413	0.595	0.488	336
	OLMo7B	0.347	0.328	0.337	336
	Qwen32B	0.405	0.504	0.449	336
	Qwen7B	0.440	0.450	0.445	336
All	LLaMa70B	0.453	0.787	0.574	2000
	LLaMa8B	0.324	0.997	0.486	2000
	ModBert	0.521	0.691	0.593	7577
	OLMo32B	0.575	0.669	0.616	2000
	OLMo7B	0.263	0.118	0.149	2000
	Qwen32B	0.601	0.713	0.648	2000
	Qwen7B	0.483	0.488	0.483	2000

Table 5: Classifier performance by dataset for professional facilitation detection. Best F1 values indicated in bold. The Qwen32B and the ModernBert models seem to be the best overall models for detecting facilitation, while the OLMo7B model completely fails.

7 Limitations

None, we are awesome.

8 Ethical Considerations

All experiments and data have passed through official ethical approval by an Ethics Committee, and are compatible with GDPR (European Parliament and Council of the European Union) and the EU AI Act (European Parliament and Council of the European Union, 2024). The annotators were compensated fairly with a monthly rate of 1000€ and signed informed consent forms before

beginning annotation and were allowed to skip uncomfortable tasks at any point, with no required justification.

In order to help research in the field of facilitation, we openly release the code that constructs [PEFK](#) under a GPLv3 license. All datasets used in [PEFK](#) are obtained from the sources provided by the respective author. We do not distribute the corpus separately, and thus no additional data license is applicable. The licenses of the original distribution apply. The Fora dataset ([Schroeder et al., 2024](#)) is not publicly available; researchers should contact the authors of Fora for access.

We use exclusively open-source models, libraries and LLMs, in order to ensure reproducibility, protect user data, and comply with relevant regulations (see [App. C.1](#)).

References

James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. 2024. [Synthetic replacements for human survey data? the perils of large language models](#). *Political Analysis*, 32(4):401–416.

Jason W. Burton, Ezequiel Lopez-Lopez, Shahar Hechtlinger, Zoe Rahwan, Samuel Aeschbach, Michiel A. Bakker, Joshua A. Becker, Aleks Berdichevskaia, Julian Berger, Levin Brinkmann, Lucie Flek, Stefan M. Herzog, Saffron Huang, Sayash Kapoor, Arvind Narayanan, Anne-Marie Nussberger, Taha Yasseri, Pietro Nickl, Abdullah Almaatouq, and 9 others. 2024. [How large language models can reshape collective intelligence](#). *Nature Human Behaviour*, 8(9):1643–1655.

Jonathan P. Chang and Cristian Danescu-Niculescu-Mizil. 2019. [Trouble on the horizon: Forecasting the derailment of online conversations as they develop](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4743–4754, Hong Kong, China. Association for Computational Linguistics.

Ming-Bin Chen, Lea Frermann, and Jey Han Lau. 2025. [WHoW: A cross-domain approach for analysing conversation moderation](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 2091–2126, Albuquerque, New Mexico. Association for Computational Linguistics.

Andy Cockburn, Philip Quinn, and Carl Gutwin. 2017. The effects of interaction sequencing on user experience and preference. *International Journal of Human-Computer Studies*, 108:89–104.

William D Crano. 1977. Primacy versus recency in retention of information and opinion change. *The Journal of Social Psychology*, 101(1):87–96.

Stefano Cresci, Amaury Trujillo, and Tiziano Fagni. 2022. [Personalized interventions for online moderation](#). In *Proceedings of the 33rd ACM Conference on Hypertext and Social Media*, HT ’22, page 248–251, New York, NY, USA. Association for Computing Machinery.

Christine De Kock, Tom Stafford, and Andreas Vlachos. 2022. [How to disagree well: Investigating the dispute tactics used on wikipedia](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3824–3837, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

Christine De Kock and Andreas Vlachos. 2021. [I beg to differ: A study of constructive disagreement in online conversations](#). In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2017–2027, Online. Association for Computational Linguistics.

European Parliament and Council of the European Union. [Regulation \(EU\) 2016/679 of the European Parliament and of the Council](#).

European Parliament and Council of the European Union. 2024. [Regulation \(eu\) 2024/1689 of the european parliament and of the council of 13 june 2024 laying down harmonised rules on artificial intelligence \(artificial intelligence act\)](#). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>. Official Journal of the European Union, L 2024/1689.

Neele Falk, Eva Vecchi, Iman Jundi, and Gabriella Lapesa. 2024. [Moderation in the wild: Investigating user-driven moderation in online discussions](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 992–1013, St. Julian’s, Malta. Association for Computational Linguistics.

Rena Gao, Ming-Bin Chen, Lea Frermann, and Jey Han Lau. 2025. [Moderation matters: Measuring conversational moderation impact in English as a second language group discussion](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 2070–2095, Vienna, Austria. Association for Computational Linguistics.

Jim Hewitt. 2003. How habitual online practices affect the development of asynchronous discussion threads. *Journal of Educational computing research*, 28(1):31–45.

Yiqing Hua, Cristian Danescu-Niculescu-Mizil, Dario Taraborelli, Nithum Thain, Jeffery Sorensen, and Lucas Dixon. 2018. [Wikiconv: A corpus of the complete conversational history of a large online collaborative community](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2818–2823, Brussels, Belgium. Association for Computational Linguistics.

Katerina Korre, Dimitris Tsirmpas, Nikos Gkoumas, Emma Cabalé, Danai Myrtzani, Theodoros Evgeniou, Ion Androutsopoulos, and John Pavlopoulos. 2025. Evaluation and facilitation of online discussions in the LLM era: A survey . In <i>Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing</i> , pages 24443–24462, Suzhou, China. Association for Computational Linguistics.	690
Anna Sophie Kümpel and Julian Unkel. 2021. (why) does comment presentation order matter for the effects of user comments? assessing the role of the availability heuristic and the bandwagon heuristic. <i>Communication Research Reports</i> , 38(4):217–228.	691
Cliff Lampe, Paul Zube, Jusil Lee, Chul Hyun Park, and Erik Johnston. 2014. Crowdsourcing civility: A natural experiment examining the effects of distributed moderation in online forums. <i>Government Information Quarterly</i> , 31(2):317–326.	692
Rajiv Movva, Pang Wei Koh, and Emma Pierson. 2024. Annotation alignment: Comparing LLM and human annotations of conversational safety . In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 9048–9062, Miami, Florida, USA. Association for Computational Linguistics.	693
Deokgun Park, Simranjit Sachar, Nicholas Diakopoulos, and Niklas Elmqvist. 2016. Supporting comment moderators in identifying high quality online news comments. In <i>Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems</i> , pages 1114–1125.	694
Joonsuk Park, Sally Klingel, Claire Cardie, Mary Newhart, Cynthia Farina, and Joan-Josep Vallbé. 2012. Facilitative moderation for online participation in erulemaking . In <i>Proceedings of the 13th Annual International Conference on Digital Government Research</i> , dg.o '12, page 173–182, New York, NY, USA. Association for Computing Machinery.	695
Perspective API. 2026. Perspective api . Accessed: 2026-02-24.	696
Luca Rossi, Katherine Harrison, and Irina Shklovski. 2024. The problems of llm-generated data in social science research . <i>Sociologica</i> , 18(2):145–168.	697
Brennan Schaffner, Arjun Nitin Bhagoji, Siyuan Cheng, Jacqueline Mei, Jay L Shen, Grace Wang, Marshini Chetty, Nick Feamster, Genevieve Lakier, and Chenhao Tan. 2024. "community guidelines make this the best party on the internet": An in-depth study of online platforms' content moderation policies . In <i>Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems</i> , CHI '24, New York, NY, USA. Association for Computing Machinery.	698
Hope Schroeder, Deb Roy, and Jad Kabbara. 2024. Fora: A corpus and framework for the study of facilitated dialogue . In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 13985–14001, Bangkok, Thailand. Association for Computational Linguistics.	699
Christopher T. Small, Ivan Vendrov, Esin Durmus, Hadjar Homaei, Elizabeth Barry, Julien Cornebise, Ted Suzman, Deep Ganguli, and Colin Megill. 2023. Opportunities and risks of llms for scalable deliberation with polis . <i>ArXiv</i> , abs/2306.11932.	700
Amaury Trujillo and Stefano Cresci. 2022. Make reddit great again: Assessing community effects of moderation interventions on r/the_donald . <i>Proc. ACM Hum.-Comput. Interact.</i> , 6(CSCW2).	701
Dimitris Tsirmpas, Ion Androutsopoulos, and John Pavlopoulos. 2026. Designing synthetic discussion generation systems: A case study for online facilitation . <i>Preprint</i> , arXiv:2503.16505.	702
Benjamin Warner, Antoine Chaffin, Benjamin Clavié, Orion Weller, Oskar Hallström, Said Taghadouini, Alexis Gallagher, Raja Biswas, Faisal Ladhak, Tom Aarsen, Griffin Thomas Adams, Jeremy Howard, and Iacopo Poli. 2025. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 2526–2547, Vienna, Austria. Association for Computational Linguistics.	703
Jiaqing Yuan and Munindar P. Singh. 2023. Conversation modeling to predict derailment . <i>Proceedings of the International AAAI Conference on Web and Social Media</i> , 17(1):926–935.	704
Justine Zhang, Ravi Kumar, Sujith Ravi, and Cristian Danescu-Niculescu-Mizil. 2016. Conversational flow in Oxford-style debates . In <i>Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 136–141, San Diego, California. Association for Computational Linguistics.	705
A Acronyms	726
NLP Natural Language Processing	727
ML Machine Learning	728
RQ Research Question	729
PEFK Prosocial and Effective Facilitation in Konversations	730
DL Deep Learning	732
IAA Inter-Annotator Agreement	733
UMOD User Moderation	734
IQ2 Intelligence Squared	735
CeRI Cornell eRulemaking Initiative	736
WHoW Why-How-Who	737
WT WikiTactics	738

Name	Type	Description
conv_id	string	The discussion’s ID. Comments under the same discussion refer to the same discussion ID.
message_id	string	The message’s (comment’s) unique ID.
reply_to	string	The ID of the comment which the current comment responds to. nan if the comment does not respond to another comment (e.g., it’s the Original Post (OP)).
user	string	Username or hash of the user that posted the comment
is_moderator	bool	Whether the user is a facilitator.
moderation_supported	bool	True if the moderation labels are directly computed from the original dataset
escalated	bool	A discussion-level measure denoting discussions which have been de-railed
escalation_supported	bool	True if the escalation labels are directly computed from the original dataset
text	string	The contents of the comment
dataset	string	The dataset from which this comment originated from
notes	JSON	A dictionary holding notable dataset-specific information
toxicity	float	The “toxicity” score given to the comment by the Perspective API
severe_toxicity	float	The “severe toxicity” score given to the comment by the Perspective API

Table 6: The full set of columns used for each comment in the [PEFK](#) dataset.

B Dataset details

B.1 Preprocessing

General

- We exclude comments with no text.
- We exclude discussions with less than two distinct participants.
- We exclude discussions which are common between wikitactics and wikiconv as well as wikidisputes and wikiconv.
 - There may be duplicate discussions between wikidisputes and wikitactics, but we allow them since they feature complementary information.

Wikiconv

- The Wikiconv corpus does not contain information about which user is a moderator/facilitator. Therefore, all comments relating to Wikiconv are tagged as non-moderators.
- Additionally, we follow the instructions of the original researchers ([Hua et al., 2018](#)), and select only discussions which have at least two comments by different users.
 - Wikipedia (thankfully) does not track users who log in with only an IP address (in the original dataset, their user_id is always set to 0 and their username is of the form 211.111.111.xxx. We consider each such username to be a separate user.
 - Due to the size of the dataset, we have to partially load it during preprocessing. Thus, there is a small chance every 100,000 records that a discussion is marked as a false negative and a part of it gets discarded.

- We only include English comments in the final dataset. We use a small, efficient library (py3langid) for language recognition, due to the large size of Wikiconv. Non-English comments are discarded before selecting valid discussions (see point above).

Wikitactics

- We infer facilitative actions by whether the comment belongs in any of the following categories:
 - Asking questions
 - Coordinating edits
 - Providing clarification
 - Suggesting a compromise
 - Contextualisation
- The above tactics are a subset of the Coordinative labels used in the WikiTactics paper. They were selected because they are not used necessarily on 1-1 discussions; they could reasonably be applied by third-party participants. Contrast them with other Coordinative labels such as “Conceding/recanting” and “I don’t know”.

Wikidisputes

- Since only 0.03% of the comments in the dataset are made by moderators, we mark the dataset as not supporting moderation.

UMOD

- Facilitative actions are marked as a gradient from 0 (no facilitation) to 1 (full facilitation). We adopt a threshold of 0.75 to consider an action as facilitative, with more than 50% annotator agreement (measured as entropy in the original dataset).

CMV-AWRY2




Figure 5: Heatmap featuring the pairwise Cohen’s κ of the annotators.

- We mark a discussion as escalated when the derailment value (from the official dataset) is in the 60th upper percentile.
- We remove deleted (“[deleted]”) comments.

B.2 Schema

Table 6 details the information available for each comment in the dataset. Some information, such as whether the person talking is a moderator / facilitator, is available only on some sub-datasets. Also note that the `is_moderator` column may hold different definitions of what we consider as moderator / facilitator. UMOD for instance considers facilitative comments made by users, while others such as Fora consider a comment facilitative if it was made by a dedicated facilitator.

B.3 Exploratory Data Analysis

We explore some initial frequencies with regard to facilitator comments in PEFK, comparing its spoken vs. written subsets, as well as the agreement and facilitator intervention frequencies in the re-annotated subset. We examine our annotated dataset in terms of class distributions and inter-annotator agreement to identify any notable patterns or outliers. Figure 3 presents the distribution of all classes across annotators. Discrepancies are evident: A11 appears particularly cautious, rarely assigning any type of intervention, while A2 is the only annotator who frequently uses negative reinforcement. Annotator A1 proved problematic, as they labeled the majority of instances as None, making them an outlier; consequently, A1 is excluded from subsequent analyses. A11 is also considered an outlier according to Figure 5

Spoken vs. written datasets in PEFK Figure 7 shows some of the structural differences of spoken and written datasets. Spoken datasets feature a significantly larger number of comments per discussion, since instances where the participants talk over each other, or interrupt the other are transcribed as separate comments. Additionally, spoken datasets feature a large amount of comments with very few words, since a lot of utterances are short but necessary for social co-ordination (e.g., “Thank you”, “Yes”, “I agree”). Figure 2 shows the




Figure 6: Number of comments plotted to the number of annotators that believe an intervention should occur.

percentage of comments made by facilitators for each of our datasets. In almost all datasets, a third of the comments are made by facilitators.⁵

Vocabulary discrepancies in PEFK datasets. We also filter the data to include only messages written by moderators, then cleaned the text by removing URLs, punctuation, and other noise. After preprocessing, we used a TF-IDF vectorization approach with n-grams (bi-grams and trigrams) to extract phrases. Crucially, the TF-IDF model was trained on all datasets combined, so that common, generic phrases were downweighted while phrases that are more distinctive to each dataset were emphasized. We then computed the average TF-IDF score of each phrase within each dataset and selected the top-ranking ones as representative expressions of moderator language in that context. The results can be found in Figure 8.

The most important phrases in CeRI suggest a formal, policy-oriented environment with an emphasis on structured participation and regulatory discussion (e.g., references to “proposed rule” and “welcome regulations”). In *fora*, the language is more conversational and facilitative, with phrases such as “great thank”, “feel free”, and “awesome thank” indicating a more informal moderation style. The Intelligence Squared (IQ2) dataset reflects the debate setting, with phrases like “ladies gentlemen” and “arguing motion”. In User Moderation (UMOD), moderators engage more in argumentation and clarification, as seen in phrases like “change view”, “missing point” reminiscing more facilitative settings and suggesting a focus on reasoning and discussion quality. The Why-How-Who (WHoW) dataset appears similar to moderated discussions or panels, with phrases like “question sir” and “answer question” indicating turn-taking and audience interaction. Finally, *WikiTactics* shows a highly procedural and rule-based style, with phrases such as “reliable sources”, “original research” and “talk page” reflecting the norms and policies of collaborative editing environments like Wikipedia.

⁵We do not include the UMOD dataset, since there is no established ground-truth (see App. B.1).




Figure 7: **Left:** Comments per discussion on the PEFK dataset. Note the significant increase in comments made in transcribed datasets such as Fora and IQ2. This deviation is explained by the fact that a single human utterance is regularly broken up to many comments in order to simulate interruptions or participants talking over each other. **Right:** Number of words per comment in all discussions in PEFK, grouped by original dataset. Comments across written datasets follow the same distribution, while spoken datasets feature a significant number of comments with very few words.




Figure 8: Most common moderator bigrams in PEFK per dataset. Function words are excluded.

C Models

C.1 LLMs

We use only open-source models, as proprietary models such as OpenAI’s GPT-5 yield much less reproducible results (Bisbee et al., 2024). Furthermore, the use of open-source models requires much fewer resources, enabling research on this task with a less demanding budget, as well as scalable deployment in real-world settings (Tsirmpas et al., 2026). It is also worth noting that since a part of our dataset (Fora) is not licensed, we can not risk leaking the data contained within to proprietary models. Generally, the use of local models protects the personal data of the users contained in our datasets, ensures that our research can be replicated in domains where anonymity must be maintained, and renders our work compatible with existing legislation (European Parliament and Council of the European Union, 2024; European Parliament and Council of the European Union).

Table 7 shows details on the LLMs used for the annotation in this study. We specifically use 4-bit quantized

versions from the unsloth huggingface repository.

Family	Version	#Params
OLMo	2 (0325)	32B
Qwen	2.5	32B
LLaMa	3.3	70B
OLMo	3	7B
Qwen	2.5	7B
LLaMa	3.1	8B

Table 7: LLMs used for annotation in this study.

C.2 Transformers

C.2.1 Training details

We construct our training and evaluation data from a unified corpus of threaded discussions, where each instance corresponds to a target comment augmented with its conversational context. For every target comment, we retrieve up to $K = 3$ preceding comments along the reply chain. Each comment is truncated to a maximum of 3000 characters to control sequence length. The input is formatted using XML-style tags, where context comments are wrapped with <CTX> and the target comment with <TGT>. The final sequence is formed by concatenating context comments in chronological order followed by the target comment. We consider binary classification tasks (e.g., is_moderator, should_intervene), with labels cast to floating point values.

We fine-tune a pretrained transformer encoder (ModernModBert-large) with an added classification head that outputs a single logit. The task is formulated as a single-label instance of multi-label classification, enabling the use of a sigmoid activation.

Training is conducted using the Hugging Face Trainer framework with several modifications. We optimize the model using binary cross-entropy with log-

its:

$$\mathcal{L} = \text{BCEWithLogitsLoss}(\mathbf{z}, \mathbf{y}; \text{pos_weight}),$$

where $\mathbf{z}$ are the predicted logits and $\mathbf{y}$ are the ground-truth labels. To account for class imbalance, the positive class is weighted as

$$\text{pos_weight} = \frac{N_{\text{negative}}}{N_{\text{positive}}}.$$

We optionally freeze the base transformer parameters and train only the classification head; this setting is enabled in our experiments to reduce computational cost and limit overfitting.

To improve computational efficiency, we employ a bucketed batching strategy that groups instances of similar length (approximated in characters) into batches of size 64. This reduces padding and improves throughput. Batches are shuffled at the bucket level at each epoch.

Training is performed for up to 80 epochs, with evaluation, logging, and checkpointing conducted once per epoch. The best model is selected based on validation loss. We apply early stopping with a warm-up phase: monitoring begins after one evaluation step, and training stops if validation loss does not improve by at least $\delta = 0.001$ for 6 consecutive evaluation steps. Model performance is evaluated using accuracy and binary F1 scores.

All experiments are implemented in PyTorch using the Hugging Face Transformers library. Data loading uses 4 worker processes, and GPU acceleration is utilized when available. We fix the random seed to 42 to ensure reproducibility.

C.2.2 One model to rule them all?

Since comments in the written and spoken datasets are fundamentally different (§2.2), we would be tempted to train different models for each. This decision is not obvious however, since it is likely that the model can internally distinguish between spoken and written speech. If that occurs, the differences between the texts should not negatively affect the model; in fact, training the model on all datasets could allow it to transfer knowledge between the two text modalities. Table 8 shows the performance of the same model when trained only on spoken and written datasets, compared to when trained on the all datasets at once. We find that the models perform better when trained on different splits.

C.2.3 Decision thresholds

An advantage of using DL-based classifiers instead of LLMs is that we can quantitatively alter the performance of the classifier by tuning the thresholds required for a comment to trigger an intervention. By making the model more conservative (increasing the threshold), we could hope that we can avoid unnecessary interventions, while making it less conservative (decreasing the threshold) we could make it more proactive.

Fig. 9 shows that the classifier has an inherently low precision score which is largely invariant to the selected

Split	Dataset	Precision	Recall	F1
Written	CeRI	0.320	0.490	0.387
	Wikitactics	0.407	0.717	0.520
Spoken	Fora	0.409	0.631	0.497
	IQ2	0.399	0.634	0.490
	WHow	0.418	0.626	0.502
All	CeRI	0.292	0.172	0.216
	Fora	0.410	0.664	0.507
	IQ2	0.406	0.623	0.491
	WHow	0.422	0.642	0.509
	Wikitactics	0.456	0.364	0.405

Table 8: Summary of classifier performance depending on the training dataset splits on the facilitator prediction task. We compare models trained on all available data with models trained separately on written and spoken datasets. Performance is generally similar.

threshold. On the other hand, recall falls off rapidly, which incentivises us to use a more liberal approach.

D Human annotation procedure

D.1 Facilitative tactics

Before the main annotation task on facilitative intervention, all annotators were extensively trained on currently established facilitative tactics. This task aimed to familiarize the annotators with what constitutes a facilitative intervention, as well as how professional facilitators operate in real-life scenarios.

Specifically, we split the annotators into three groups, each of which was assigned a different taxonomy of facilitative interventions (Chen et al., 2025; Schroeder et al., 2024; Park et al., 2012). The annotators were tasked with classifying 830 facilitated comments from a sample of continuous discussions in PEFK. The guidelines and results of this annotation are not shown in this study, but they are available upon request.

D.2 Expert facilitators vs. expert annotators

In designing our data annotation procedure, we intentionally prioritized the recruitment of expert annotators over expert moderators. This distinction is pivotal to our research design: whereas expert moderators typically engage in post hoc intervention—addressing content only after its publication—expert annotators facilitate a more proactive analysis of the discourse as it unfolds. However, acknowledging the importance of moderation experience we also include a professional moderator into the workflow to provide a more just comparison. In other words, we operate under the premise that the majority of discussion facilitation and moderation is conducted organically by the participants themselves through user-driven flagging mechanisms or actual dialogue.

All does not exhibit the characteristics one would expect from a distinct “expert” annotator, but instead




Figure 9: Classifier performance (**left**: written datasets, **right**: spoken datasets) when tuning the decision threshold on the test-set. The higher the threshold, the more conservative the classifier is encouraged to become when predicting a facilitator intervention.

behaves as a highly balanced and representative member of the group. In contrast to other annotators who show clear directional biases—either favoring positive reinforcement (e.g., A1, A7, A8) or leaning negative (e.g., A2, A6)—A11 distributes judgments almost evenly between positive (28.80%) and negative (28.12%) labels, resulting in a near-zero difference (0.69 percentage points). At the same time, A11’s intervention rate (56.92%) is almost identical to the group average, indicating that it is neither more strict nor more lenient in deciding when to intervene. When considered alongside the moderate correlation scores with other annotators (mean 0.38), this pattern suggests that A11 does not consistently align more strongly with the group than others do, nor does it systematically deviate in a principled way. Instead, A11 appears to occupy the “center” of the annotator space: its behavior closely reflects the aggregate tendencies of the group without amplifying any particular bias. Consequently, the data do not support interpreting A11 as a clear expert or gold-standard annotator; rather, A11 is better characterized as a calibrated, balanced annotator whose decisions approximate the overall distribution of judgments in the dataset.

D.3 Positive vs. negative reinforcement

The distinction between positive/negative reinforcement is present to suggested facilitative tactics taxonomies (Table 2), including prior work such as Fora (Schroeder et al., 2024) WHoW (Chen et al., 2025), and CeRI (Park et al., 2012). We conservatively estimate whether the tactics described in each taxonomy correspond to our definitions of positive and negative reinforcement.

The facilitative taxonomy proposed in Fora does not include any tactics that could be unambiguously considered as negative reinforcement, which is expected, since facilitators deployed in the relevant study were supposed to only act as coordinators. Notably, several tactics may simultaneously exhibit both positive and negative characteristics across all taxonomies, depending on context and the exact implementation used by the facilitator. We therefore use these taxonomies only as

guidance during annotation, and refer the reader to the original works for a more comprehensive description and analysis of facilitation strategies.

D.4 Label Transformation

To efficiently calculate inter- and intra-annotator agreement, we transform the fine-grained, ordinal label per type adopting a multiclass schema. We define a threshold of ≥ 3 and then using Table ?? we relabel the items.

D.5 Intra-annotator agreement

To evaluate the consistency of individual annotators, we measured intra-annotator agreement using purposefully repeated instances of identical discussions.

Agreement was computed as a strict match between the two annotations of the same instance (duplicated pair). For each duplicated pair, agreement was assigned a value of 1 if both labels were identical and 0 otherwise in order to calculate a simple pairwise agreement. Hence, the final intra-annotator agreement for each annotator was then calculated as the mean of these binary agreement scores.

We report results in Table 13. Most annotators demonstrated near-perfect consistency, with several achieving 100% agreement. However, annotator A1 exhibited substantially lower consistency (55.56%). This is due to the fact that he did not manage to complete all annotation tasks. Annotator A3 also showed reduced stability (66.67%), suggesting potential variability in interpretation or labeling behavior.

E Annotation guidelines

E.1 Human guidelines

Annotation Guidelines: Type of Reinforcement Needed in Comment Threads

Purpose of the Task

These guidelines explain how to annotate comments in conversation threads using three parallel rating columns:

- Positive Reinforcement (1–5)
- Negative Reinforcement (1–5)

Taxonomy & Tactic	Type
Fora	
Express appreciation	+
Express agreement or affirmation	+
Open invitation to participants	
Specific invitation to participant	
Model examples	
Follow-up question	
Make connections	
WHoW	
Informational	
Coordinative	-
Social	+
Probing	
Confronting	
Instruction	-
Interpretation	
Supplement	
CeRI	
Welcoming	+
Encouragement / appreciation	+
Maintaining civil discourse	-
Explaining why comment not removed	
Indicating irrelevant/off-point comments	-
Pointing to relevant information	
Pointing out effective commenting	+
Asking for more information/sources	
Asking for solutions/alternatives	
Encourage engagement with others	
Pose question to community	

Table 9: Facilitation taxonomies (bold) with associated tactics (normal text) and reinforcement type. +: positive reinforcement, -: negative reinforcement, empty if the tactic does not clearly correspond to any category.

Annotator	Correlation with A11
A6	0.492
A10	0.464
A2	0.432
A8	0.430
A1	0.396
A9	0.385
A4	0.383
A3	0.363
A12	0.301
A7	0.279
A5	0.249
Mean	0.380
Min	0.249
Max	0.492

Table 10: Correlation of Expert Moderators with Other Annotators

- No Reinforcement Needed (1–5)

Your job is to read each comment in its conversational context and assign a strength/certainty rating (1–5) for each category, reflecting how that category applies. All fields need to be filled and the categories are not necessarily mutually exclusive.

Even though your intervention will be delivered at

the end of the thread, the reinforcement can refer to any point in the conversation, not just at the end. It can refer to something said earlier, not just the last comment.

You are not being asked whether that conversation correlates to positive or negative reinforcement, but rather whether you, personally, would have used positive or negative reinforcement at a given point. Keep in mind that too many interventions on the moderator's part may result in an adverse reaction from the participants; no one wants to be constantly told what to say or how to say it.

Important Notes

- You must add the rationale behind your response in the rational/comments field when:
 - your positive score > 3, or
 - your negative score > 3, or
 - you encounter a particular case (e.g., an instance which could fit both negative and positive reinforcement categories).
- If you encounter data that is malformed or noisy to the point that it renders comprehension difficult, flag the instance using the malformed field.

What Each Label Means

Positive Reinforcement

Positive reinforcement is defined as an intervention whose purpose is to encourage a specific behavior. For example, a moderator may ask people to participate or ask a person to elaborate. Positive reinforcement is not related to positive speech; the moderator may be strict or even rude when attempting to encourage specific behaviors (e.g., "I will not tolerate any further discussion on this topic. It's been decided and we're moving on.").

Rate this when a conversation demonstrates behavior that should be encouraged or reinforced in a synchronous exchange. Your task is not simply to identify positive behaviors, but to identify moments where reinforcing them would help sustain or build constructive conversational momentum.

This includes, but is not limited to:

- Helpful or clarifying explanations
- Respectful disagreement
- Constructive feedback
- Polite, empathetic, or supportive tone
- Efforts to de-escalate conflict
- Any contribution that meaningfully improves the quality or direction of the conversation

Negative Reinforcement

The purpose of negative reinforcement is to discourage unwanted or disruptive behaviors. This is correlated, but not necessarily closely related to threats or enforcement of disciplinary action. Negative reinforcement is not related to negative speech; the moderator may be kind and diplomatic when discouraging behaviors (e.g., "I've received multiple complaints about your behavior on this forum. Please refrain from making personal attacks and inflammatory comments.").

Rate this when the comment reflects harmful, disruptive, or inappropriate behavior that should be discour-

Annotator	Negative (%)	No Reinforcement (%)	Positive (%)	Intervention Rate (%)	Pos - Neg (pp)
A1	27.61	31.99	40.40	68.01	12.78
A2	27.65	49.62	22.73	50.38	-4.92
A3	39.77	13.65	46.57	86.35	6.80
A4	23.82	47.59	28.59	52.41	4.77
A5	20.78	47.95	31.27	52.05	10.50
A6	24.73	53.11	22.16	46.89	-2.57
A7	19.84	46.21	33.95	53.79	14.11
A8	24.47	38.93	36.60	61.07	12.13
A9	25.52	43.51	30.97	56.49	5.45
A10	24.52	43.32	32.15	56.68	7.63
A11	28.12	43.08	28.80	56.92	0.69
A12	17.81	65.40	16.79	34.60	-1.01

Table 11: Behavioral Summary of Annotators (Percentages)

Positive	Negative	No Reinforcement	Final
✗	✗	✗	No (0)
✓	✗	✗	Positive (1)
✗	✓	✗	Negative (2)
✗	✗	✓	No (0)
✓	✓	✗	Both (3)
✓	✗	✓	Positive (1)
✗	✓	✓	Negative (2)
✓	✓	✓	Both (3)

Table 12: Reinforcement label transformations.

Annotator	Percentage Agreement (%)
A1	55.56
A2	100.00
A3	66.67
A4	88.89
A5	100.00
A6	100.00
A7	100.00
A8	100.00
A9	100.00
A10	100.00
A11	100.00

Table 13: Intra-annotator percentage agreement.

aged, as well as cases where some participants dominate the conversation, or when the conversation starts to go off topic. Other examples include:

- Hostile, dismissive, or disrespectful language
- Personal attacks or insults
- Repeated interruptions or monopolizing the exchange
- Escalating conflict rather than resolving it
- Contributions that reduce the overall quality of the conversation

No Reinforcement Needed

Rate this when the comment is neutral, unremarkable, or does not affect the tone or direction of the conversation. This includes:

- Basic factual statements
- Clarifying questions

- Minimal or vague comments
- Harmless but non-notable contributions

How to Make Annotation Decisions

Look at Intent and Impact

Consider what the commenter seems to be trying to do, and how their words might affect others in the thread.

Use the Context

Always check the surrounding messages. A comment that looks fine on its own might escalate tension when viewed in context—and vice versa.

Judge Behavior, Not Opinions

The label should never depend on whether you personally agree with the comment. Focus on tone, clarity, constructiveness, and how it influences the conversation.

Remain Consistent

Apply the same standards across all comments. Avoid guessing or over-interpreting; use what is present in the text.

E.2 LLM prompts

E.2.1 Yes/no intervention

You are an expert annotator specializing in online moderation and facilitation. You will be presented with a comment and the comments made before it. Choose whether you decide to intervene at this point in the discussion.

Key Rules for Annotating:

Look at Intent and Impact Consider what the commenter seems to be trying to do, and how their words might affect others in the thread.

Use the Context Always check the surrounding messages. A comment that looks fine on its own might escalate tension when viewed in context—and vice versa.

Judge Behavior, Not Opinions The label should never depend on whether you personally agree with the comment. Focus on tone, clarity, constructiveness, and how it influences the conversation.

Remain Consistent Apply the same standards across all comments. Avoid guessing or over-interpreting; use what is present in the text.

Output Respond whether you would intervene in this discussion as a facilitator on a scale of 1-5, where 1 is

1221	"I'm sure I don't need to intervene", and 5 is "I'm sure		
1222	I need to intervene"		
1223	E.2.2 Positive vs. negative reinforcement		
1224	You are an expert facilitator. You will be presented		
1225	with a discussion. Decide whether you would write		
1226	a comment intervening after the last comment in the		
1227	discussion.		
1228	The Core Rating Task		
1229	For every comment, you must write out a response of		
1230	1 (uncertain) to 5 (certain):		
1231	Positive Reinforcement		
1232	Defined as an intervention whose purpose is to en-		
1233	courage a specific behavior. For example, a moderator		
1234	may ask people to participate, ask a person to elaborate		
1235	etc.		
1236	Positive reinforcement is not related to positive		
1237	speech; the moderator may be strict or even rude when		
1238	attempting to encourage specific behaviors (e.g., "I will		
1239	not tolerate any further discussion on this topic. It's		
1240	been decided and we're moving on.").		
1241	Rate this when a conversation demonstrates behav-		
1242	ior that should be encouraged or reinforced in a syn-		
1243	chronous exchange. Your task is not simply to identify		
1244	positive behaviors, but to identify moments where rein-		
1245	forcing them would help sustain or build constructive		
1246	conversational momentum.		
1247	This includes, but is not limited to:		
1248	• Helpful or clarifying explanations		
1249	• Respectful disagreement		
1250	• Constructive feedback		
1251	• Polite, empathetic, or supportive tone		
1252	• Efforts to de-escalate conflict		
1253	• Any contribution that meaningfully improves the		
1254	quality or direction of the conversation		
1255	Negative Reinforcement		
1256	The purpose of negative reinforcement is to discour-		
1257	age unwanted or disruptive behaviors. This is correlated,		
1258	but not necessarily closely related to threats or enforce-		
1259	ment of disciplinary action.		
1260	Negative reinforcement is not related to negative		
1261	speech; the moderator may be kind and diplomatic when		
1262	discouraging behaviors (e.g., "I've received multiple		
1263	complaints about your behavior on this forum. Please		
1264	refrain from making personal attacks and inflammatory		
1265	comments.").		
1266	Rate this when the comment reflects harmful, disrupt-		
1267	ive, or inappropriate behavior that should be discour-		
1268	aged, as well as cases where some participants dominate		
1269	the conversation, or when the conversation starts to go		
1270	off topic.		
1271	Other examples include:		
1272	• Hostile, dismissive, or disrespectful language		
1273	• Personal attacks or insults		
1274	• Repeated interruptions or monopolizing the ex-		
1275	change		
1276	• Escalating conflict rather than resolving it		
1277	• Contributions that reduce the overall quality of the		
1278	conversation		
	No Reinforcement		1279
	Rate this when the comment is neutral, unremarkable,		1280
	or does not affect the tone or direction of the conversa-		1281
	tion.		1282
	This includes:		1283
	• Basic factual statements		1284
	• Clarifying questions		1285
	• Minimal or vague comments		1286
	• Harmless but non-notable contributions		1287
	Key Rules for Annotating		1288
	Look at Intent and Impact		1289
	Consider what the commenter seems to be trying to do,		1290
	and how their words might affect others in the thread.		1291
	Use the Context		1292
	Always check the surrounding messages. A comment		1293
	that looks fine on its own might escalate tension when		1294
	viewed in context—and vice versa.		1295
	Judge Behavior, Not Opinions		1296
	The label should never depend on whether you person-		1297
	ally agree with the comment. Focus on tone, clarity,		1298
	constructiveness, and how it influences the conversa-		1299
	tion.		1300
	Remain Consistent		1301
	Apply the same standards across all comments. Avoid		1302
	guessing or over-interpreting; use what is present in the		1303
	text.		1304
	Output		1305
	Respond whether you would intervene in this discus-		1306
	sion as a facilitator.		1307
	Always respond with the following format and noth-		1308
	ing else:		1309
	Positive: 1-5		1310
	Negative: 1-5		1311
	No Reinforcement: 1-5		1312
	E.2.3 Single-comment facilitation prediction		1313
	You are an expert moderator/facilitator. You will be		1314
	presented with a comment (TGT) and the comments		1315
	made before it (CTX, one for each comment in turn).		1316
	Choose whether you decide to intervene at this point in		1317
	the discussion.		1318
	Key Rules for Annotating:		1319
	Look at Intent and Impact Consider what the com-		1320
	menter seems to be trying to do, and how their words		1321
	might affect others in the thread.		1322
	Use the Context Always check the surrounding mes-		1323
	sages. A comment that looks fine on its own might es-		1324
	calate tension when viewed in context—and vice versa.		1325
	Judge Behavior, Not Opinions The label should		1326
	never depend on whether you personally agree with the		1327
	comment. Focus on tone, clarity, constructiveness, and		1328
	how it influences the conversation.		1329
	Remain Consistent Apply the same standards across		1330
	all comments. Avoid guessing or over-interpreting; use		1331
	what is present in the text.		1332
	Output Respond whether you would intervene in this		1333
	discussion as a facilitator on a scale of 1-5, where 1 is		1334
	"I'm sure I don't need to intervene", and 5 is "I'm sure I		1335

need to intervene". Answer with a single number, and nothing else.

E.2.4 Single-comment facilitation detection

You are an expert moderator/facilitator. You will be presented with a comment (TGT) and the comments made before it (CTX, one for each comment in turn). Choose whether the comment you are reading (TGT) is a comment made by a moderator/facilitator.

Key Rules for Annotating:

Look at Intent and Impact Consider what the commenter seems to be trying to do, and how their words might affect others in the thread.

Use the Context Always check the surrounding messages. A comment that looks fine on its own might escalate tension when viewed in context—and vice versa.

Judge Behavior, Not Opinions The label should never depend on whether you personally agree with the comment. Focus on tone, clarity, constructiveness, and how it influences the conversation.

Remain Consistent Apply the same standards across all comments. Avoid guessing or over-interpreting; use what is present in the text.

Output Respond whether the comment you are reading (TGT) is a comment made by a moderator/facilitator on a scale of 1-5, where 1 is "Im sure it was not made by a moderator/facilitator", and 5 is "I'm sure it was made by a moderator/facilitator." Answer with a single number, and nothing else.