

To intervene or not to intervene

Anonymous ACL submission

Abstract

This document is a supplement to the general instructions for *ACL authors. It contains instructions for using the $\text{\LaTeX}$ style files for ACL conferences. The document itself conforms to its own specifications, and is therefore an example of what your manuscript should look like. These instructions should be used both for papers submitted for review and for final versions of accepted papers.

1 Introduction

Dimitris, Katerina

2 Related Work

2.1 Human facilitators

Katerina, Dimitris Add research on how humans typically behave when facilitating.

2.2 Facilitation detection

Dimitris Add work that already tried to answer this question (Fora, WHoW, the second whow aussie paper).

3 Dataset

3.1 Description

Dimitris. Transfer basic information from project github to here including stats, tables and licensing.

3.2 Text-based vs. transcribed datasets

Dimitris, Katerina Our datasets are fundamentally different depending on whether the data are transcripts or actual text data. Dimitris will add technical examples (i.e., how the transcriptions have a different format, e.g., comments that cut each other off). Katerina will add social science research on how oral and written speech differs. Thus, we run separate experiments for each of the two dataset types.

4 How to decide when to intervene?

4.1 Do not decide

Dimitris. What most platforms do is to use flagging and reports to decide when to intervene. Mention Grok as a facilitative agent that works by invocation only. Briefly talk about advantages and disadvantages of current approach.

4.2 Use humans

4.2.1 Merits and disadvantages

Katerina. Humans may be the “best” classifier since they understand context and social behavior best, but they are not scalable (cite our survey). Dimitris can add cost calculations for a hypothetical scenario to drive point home.

4.2.2 Experimental setup

Katerina. Annotation study. We leave experimental details for the appendix, but mention expert annotators, basic procedure, data and setup here. Also talk about positive vs negative reinforcement.

4.2.3 Results

Katerina. Inter-annotator agreement and discussion.

4.3 Use LLMs

4.3.1 Can we detect facilitation interventions?

Dimitris. We use LLMs and transformer-based models to decide whether the person talking is a facilitator.

4.3.2 Can LLMs decide when to intervene?

Dimitris, Katerina. No. We show this with two experiments. One is agreement between humans and LLMs on the annotation task (strong baseline). The other is by predicting when a facilitator would speak (weak baseline).

5 Conclusion

Humans are good but expensive. Our research is good. We wasted a good amount of money on this, please accept this.