

Explanation of Revisions

We thank the reviewers for their insight, engagement and recommendations for our work. Below is a list of all major changes from the previous submission cycle:

1. A formatting oversight on our part led to the accidental adoption of 10pt fontsize instead of the ACL standard of 11pt. The error has been corrected.
2. §5.3 from the previous draft has been moved to Appendix F.1, since both reviewers bRUY and hmYL understandably found it difficult to connect the experiment of facilitator detection with facilitator prediction. Since the connection between the two is supplementary to our results, we made the decision to move the section outside the main text.
3. After a discussion with reviewer hmYL we also dropped the framing that facilitator detection functions as an upper-bound estimation for facilitator prediction. The task is now clearly marked as a Natural Language Understanding task, as mentioned in the previous manuscript.
4. We added ablations both for instruction sensitivity (including prompts specifically mentioning that facilitator overreach should be treated as an important concern) and for LLM self-consistency after a recommendation of both reviewers bRUY and hmYL in Appendix F.2.
5. We added additional information about the participants of the study and the experimental procedure (including pilot studies, interviews and meetings) in Appendix E, as suggested by reviewer hmYL. We also motivated our decision to train the participants in §4.1, after the recommendations of reviewer bRUY.
6. We release the full pilot dataset including survey participants rationales as requested by reviewer hmYL (see Supplementary), including some analysis presented in Appendix E. We mention that the results from the pilot are not directly comparable to our main results, for the reasons outlined in the same Appendix. We also openly acknowledge in the Limitations that, with hindsight, we should have leaned more on these rationales while conducting the survey.

7. We added an experiment showing that LLM facilitators may be systematically influenced by discussion domain, as correctly hypothesized by reviewers bRUY and Wbzh in Appendix F.5. The results are interesting, although not conclusive enough to present them as a main finding. We include a brief discussion in the relevant Appendix section.
8. We added general clarifications for points raised by all reviewers in the main text. We decided to keep the example used in Table 1, although we explicitly note that the three definitions of facilitative interventions are not mutually exclusive in the relevant section, as suggested by reviewer hmYL.
9. We reworked Figures 2,3,4,8, to make them colorblind-friendly.